

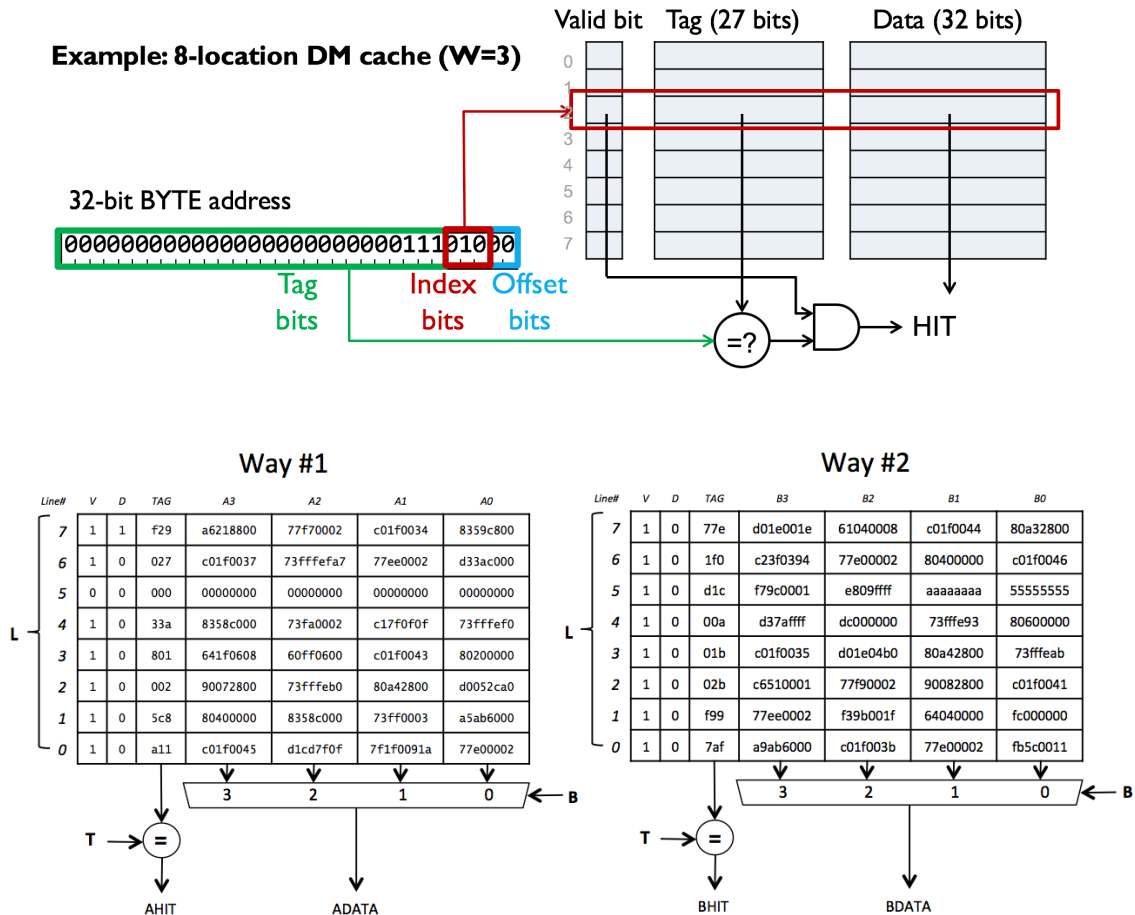
## 6.004 Recitation Problems

### L15 – Caches

Keep the most often-used data in a small, fast SRAM (often local to CPU chip). The reason this strategy works: LOCALITY.

- *Temporal locality: If a location has been accessed recently, it is likely to be accessed (reused) soon*
- *Spatial locality: If a location has been accessed recently, it is likely that nearby locations will be accessed soon*

$$\text{AMAT(Average Memory Access Time)} = \text{HitTime} + \text{MissRatio} * \text{MissPenalty}$$



**Replacement strategy choices:**

- Oracle, OPT (optimal)
- RANDOM
  - evict a random way
- FIFO (first-in, first-out)
  - give every way equal residency
- LRU (least-recently used)
  - The one used least recently will get evicted first

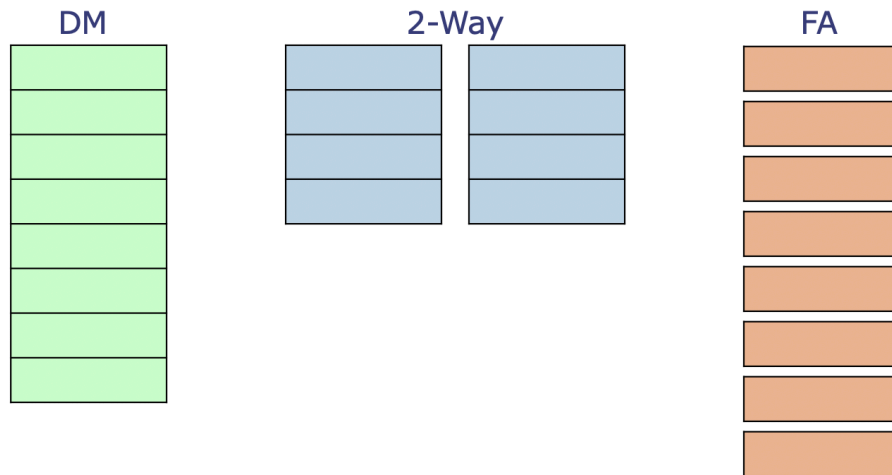
**Write-policy choices:**

- Write-through (write data go to cache and memory)
  - Update the main memory as well as the cache on every write
  - Replacing a cache entry is simple (just overwrite new block)
  - Memory write causes significant delay
- Write-back (write data only goes to the cache)
  - Only the cache entry is updated on each cache write so main memory and the cache data are inconsistent.
  - Add “dirty” bit to the cache entry to indicate whether the data in the cache entry must be committed to memory.
  - Replacing a cache entry requires writing the data back to memory before replacing the entry if it is “dirty”.

(credit: Stanford EE108b lecture slides)

## Example: Comparing Hit Rates

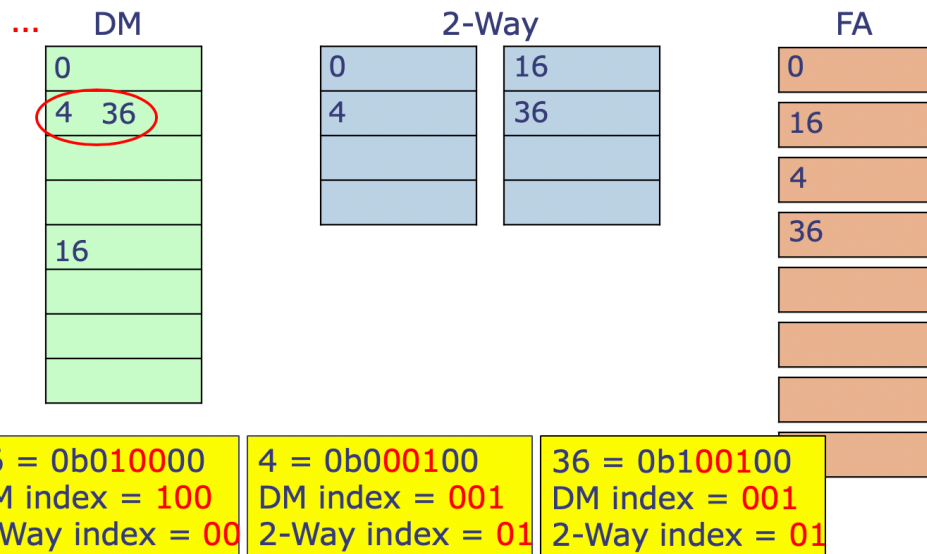
- 3 Caches: DM, 2-Way, FA: each has 8 (4 byte) words



MIT 6.004 Spring 2019

## Example: Comparing Hit Rates

Access following addresses repeatedly: 0, 16, 4, 36,



MIT 6.004 Spring 2019

## Example: Comparing Hit Rates

Access following addresses repeatedly: 0, 16, 4, 36,

DM	2-Way		FA
0	0	16	0
4, 36	4	36	16
			4
			36
16			

DM: 50% hit rate  
 2-Way: 100% hit rate  
 FA: 100% hit rate

MIT 6.004 Spring 2019

## Example 2: Comparing Hit Rates

▪ Access: 0, 4, 8, 12, 16, 20, 24, 28, 32, ...

DM	2-Way		FA
0, 32	0, 32, 16	16, 0, 32	0 32
4	4	20	4 0
8	8	24	8 4
12	12	28	12 8
16			16 12
20			20 16
24			24 20
28			28 24

DM: Hit rate = 7/9    2-Way: Hit rate = 6/9    FA: Hit rate = 0%

MIT 6.004 Spring 2019

## Example 3: Comparing Hit Rates

- Access: 0, 4, 8, 12, 32, 36, 40, 44, 16, ...

DM	2-Way		FA
0, 32	0, 16, 32	32, 0, 16	0 16
4, 36	4	36	4 0
8, 40	8	40	8 4
12, 44	12	44	12 8
16			32 12
			36 32
			40 36
			44 40

DM: Hit rate =  $1/9$     2-Way: Hit rate =  $6/9$     FA: Hit rate = 0%

MIT 6.004 Spring 2019

### Problem 1.

- (A) The timing for a particular cache is as follows: checking the cache takes 1 cycle. If there's a hit the data is returned to the CPU at the end of the first cycle. If there's a miss, it takes 10 *additional* cycles to retrieve the word from main memory, store it in the cache, and return it to the CPU. If we want an average memory access time of 1.4 cycles, what is the minimum possible value for the cache's hit ratio?

Minimum possible value of hit ratio: \_\_\_\_\_

- (B) If the cache block size, i.e., words/cache line, is doubled but the total number of data words in the cache is unchanged, how will the following cache parameters change? Please circle the best answer.

# of offset bits: UNCHANGED ... +1 ... -1 ... 2x ... 0.5x ... CAN'T TELL

# of tag bits: UNCHANGED ... +1 ... -1 ... 2x ... 0.5x ... CAN'T TELL

# of cache lines: UNCHANGED ... +1 ... -1 ... 2x ... 0.5x ... CAN'T TELL

Consider a direct-mapped cache with 64 total data words with 1 word/cache line, which uses a LRU replacement strategy and a write-back write strategy. This cache architecture is used for parts (C) through (F).

- (C) If cache line number 5 is valid and its tag field has the value 0x1234, what is the address in main memory of the data word currently residing in cache line 5?

**Main memory address of data word in cache line 5: 0x\_\_\_\_\_**

The program shown on the right repeatedly executes an inner loop that sums the 16 elements of an array that is stored starting in location 0x310.

**The program is executed for many iterations**, then a measurement of the cache statistics is made during one iteration through all the code, i.e., starting with the execution of the instruction labeled `outer_loop`: until just before the next time that instruction is executed.

```
. = 0                // tell assembler to start at
                        // address 0
outer_loop:
    addi x4, x0, 16    // initialize loop index J
    mv x1, x0          // x1 holds sum, initially 0

loop:                  // add up elements in array
    subi x4, x4, 1     // decrement index
    slli x2, x4, 2     // convert to byte offset
    lw x3, 0x310(x2)   // load value from A[J]
    add x1, x1, x3     // add to sum
    bne x4, x0, loop   // loop until all words are summed

    j outer_loop      // perform test again!
```

- (D) In total, how many instruction fetches occur during one complete iteration of the outer loop?  
How many data reads?

**Number of instruction fetches: \_\_\_\_\_**

**Number of data reads: \_\_\_\_\_**

- (E) How many instruction fetch misses occur during one complete iteration of the outer loop?  
How many data read misses? Hint: remember that the array starts at address 0x310.

**Number of instruction fetch misses: \_\_\_\_\_**

**Number of data read misses: \_\_\_\_\_**

(F) What is the hit ratio measured after one complete iteration of the outer loop?

**Hit ratio:** \_\_\_\_\_

**Problem 2.**

The RISC-V Engineering Team is working on the design of a cache. They've decided that the cache will have **a total of  $2^{10} = 1024$  data words**, but are still thinking about the other aspects of the cache architecture.

First assume the team chooses to build a direct-mapped write-back cache with a block size of 4 words.

(A) Please answer the following questions:

**Number of lines in the cache:** \_\_\_\_\_

**Number of bits in the tag field for each cache entry:** \_\_\_\_\_

(B) This cache takes *2 clock cycles* to determine if a memory access is a hit or a miss and, if it's a hit, return data to the processor. If the access is a miss, the cache takes *20 additional clock cycles* to fill the cache line and return the requested word to the processor. If the hit rate is 90%, what is the processor's average memory access time in clock cycles?

**Average memory access time assuming 90% hit rate (clock cycles):** \_\_\_\_\_

Now assume the team chooses to build a 2-way set-associative write-back cache with a block size of 4 words. *The total number of data words in the entire cache is still 1024.* The cache uses a LRU replacement strategy.

(C) Please answer the following questions:

**Address bits used as offset (including byte offset):** A[ \_\_\_\_ : \_\_\_\_ ]

**Address bits used as cache line index:** A[ \_\_\_\_ : \_\_\_\_ ]

**Address bits used for tag comparison:** A[ \_\_\_\_ : \_\_\_\_ ]

(D) To implement the LRU replacement strategy this cache requires some additional state for each set. How many state bits are required for each set?

**Number of state bits needed for each set for LRU:** \_\_\_\_\_

To test this set-associative cache, the team runs the benchmark code shown on the right. The code sums the elements of a 16-element array. The first instruction of the code is at location 0x0 and the first element of the array is at location 0x10000. Assume that the cache is empty when execution starts and remember *the cache has a block size of 4 words*.

```
. = 0x0
mv x3, x0 // index
mv x1, x0 // sum
// x4 = 0x10000
lui x4, 0x10
```

(E) How many instruction misses will occur when running the benchmark?

**Number of instruction misses when running the benchmark:** \_\_\_\_\_

```
L: add x5, x4, x3
lw x2, 0(x5)
add x1, x1, x2
addi x3, x3, 4
slti x2, x3, 64
bnez x2, L
unimp // halt
```

(F) How many data misses (i.e., misses caused by the memory access from the LD instruction) will occur when running the benchmark?

**Number of data misses when running the benchmark:** \_\_\_\_\_

```
. = 0x10000
A: .word 0x1
.word 0x2
...
.word 0xF
.word 0x10
```

(g) What's the exact hit rate when the complete benchmark is executed?

**Benchmark hit rate:** \_\_\_\_\_



### Problem 3.

Assume, the program shown on the right is being run on a RISC-V processor with a cache with the following parameters:

- **2-way set-associative**
- **block size of 2**, i.e., 2 data words are stored in each cache line
- total number of data words in the cache is **32**
- **LRU** replacement strategy

(A) The cache will divide the 32-bit address supplied by the processor into three fields: B bits of block offset (including byte offset bits), L bits of cache line index, and T bits of tag field. Based on the cache parameters given above, what are the appropriate values for B, L, and T?

value for B: \_\_\_\_\_

value for L: \_\_\_\_\_

value for T: \_\_\_\_\_

```
. = 0x240          // start of program
test:
    addi x4, x0, 16 // initialize loop index J
                      // to size of array
    mv x1, x0       // x1: sum

loop:                // add up elements in array
    subi x4, x4, 1   // decrement index
    slli x2, x4, 2    // convert to byte offset
    lw x3, 0x420(x2) // load value from A[J]
    add x1, x1, x3    // add to sum
    bnez x4, loop     // loop N times

    j test            // perform test again!

// allocate space to hold array
. = 0x420
A: .word A[0]
   .word A[1]
   ...
```

(B) If the SLLI instruction is resident in a cache line, what will be its cache line index? the value of the tag field for the cache?

Cache line index for SLLI when resident in cache: \_\_\_\_\_

Tag field for SLLI when resident in cache: 0x\_\_\_\_\_

(C) Given that the code begins at address 0x240 and the array begins at address 0x420, and that there are 16 elements in the array as shown in the code above, list *all* the values  $j$  ( $0 \leq j < 16$ ) where the location holding the value  $A[j]$  will map to the same cache line index as the SLLI instruction in the program.

List all  $j$  where  $A[j]$  have the same cache line index as SLLI: \_\_\_\_\_

(D) If the outer loop is run many times, give the steady-state hit ratio for the cache, i.e., assume that the number of compulsory misses as the cache is first filled are insignificant compared to the number of hits and misses during execution.

Steady-state hit ratio (%): \_\_\_\_\_