

Processor Pipelining

Arvind

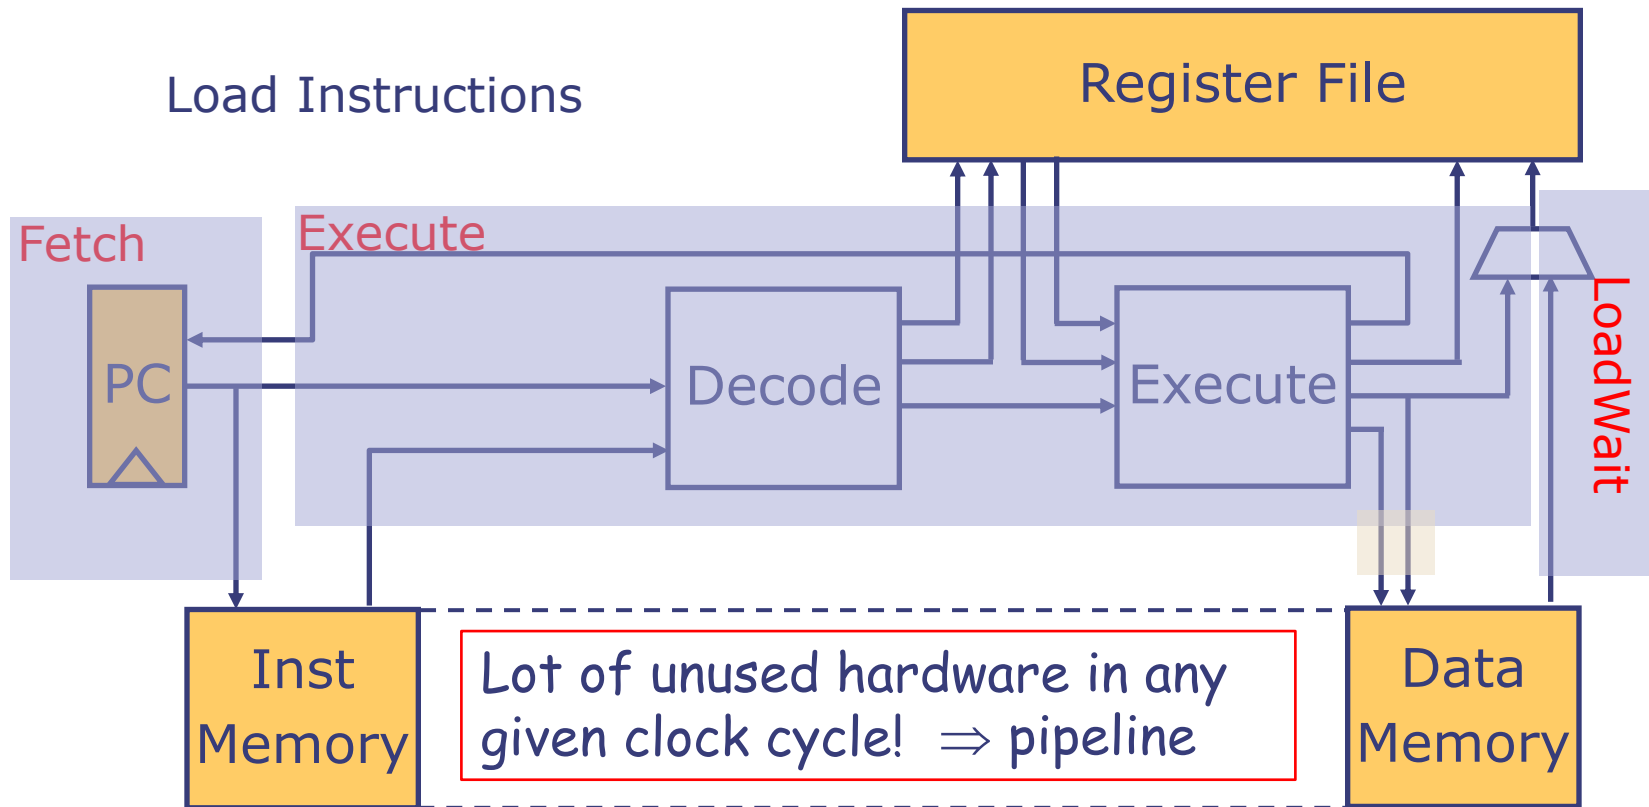
Computer Science & Artificial Intelligence Lab
M.I.T.

Processor pipelines

- Pipelining processor provides the ultimate challenge in computer architecture
 - Stringent correctness requirements
 - Requires speculative execution of instructions to pipeline at all!
 - Requires dealing with a variety of feedback
 - The goal is always to achieve highest performance but within a given area and power budget

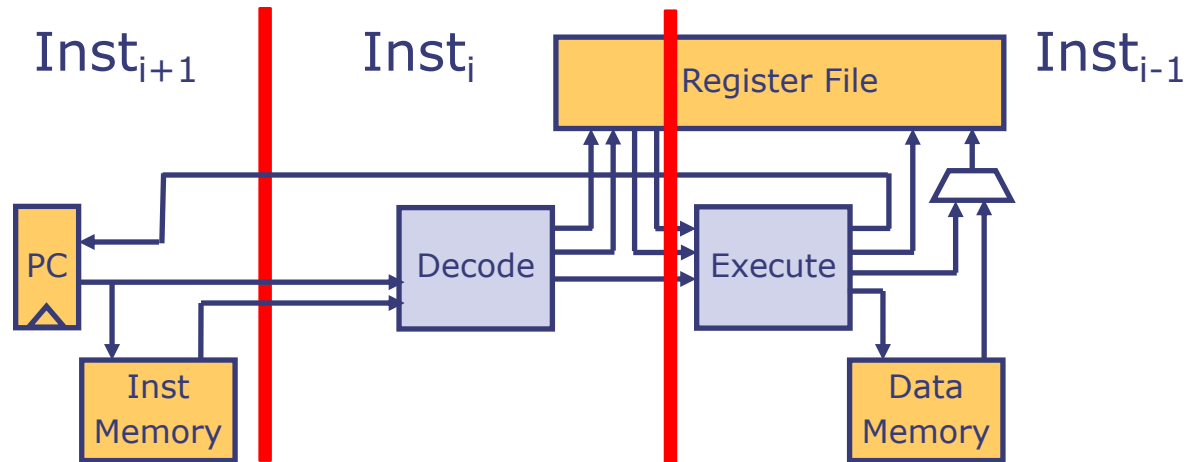
Multicycle RISC-V: *Analysis*

Lab 6



- Assuming 20% load instructions, and memory latency of one, the average number of cycles per instruction:
 - $2 \times .8 + 3 \times .2 = 2.2$ higher memory latency will make this number much worse

New problems in pipelining instructions over arithmetic pipelines




- *Control hazard*: pc for $Inst_{i+1}$ is not known until at least $Inst_i$ is decoded. So which instruction should be fetched?
 - Solution: *Speculate and squash* if the prediction is wrong
- *Data hazard*: $Inst_i$ may be dependent on $Inst_{i-1}$, and thus, it must wait for the effect of $Inst_{i-1}$ on the state of the machine (pc, rf, dMem) to take place
 - Solution: *Stall* instruction $Inst_i$ until the dependency is resolved
 - Number of stalls can be reduced by *bypassing*, that is by providing additional datapaths

Plan

1. Develop a two-stage pipeline by providing a solution for *control hazards*
2. Develop a three-stage pipeline by also providing a solution for *data hazards*

Many code fragments from the multicycle implementation (Lab 6) are reusable


Control hazard



We will offer a solution that is independent of how many cycles each stage takes


- Fetch stage initiates instruction fetches and sends them to Execute stage via f2d. It speculatively updates pc to pc+4
- Execute stage picks up instruction from f2d and executes it. It may take one or more cycles to do this
- These two stages operate independently except in case of a branch misprediction when Execute redirects the pc to the correct pc

Timing diagrams and bubbles



Multicycle Processor

Execution of
 $I_0; I_1; I_2; I_3; \dots$
 only I_2 is a load
 instruction



Two-stage Pipeline

Suppose I_0 is a
 branch instruction
 which jumps to I_7
 instead of I_1

If squashing takes more than
 one cycle then I_7 will get
 further delayed

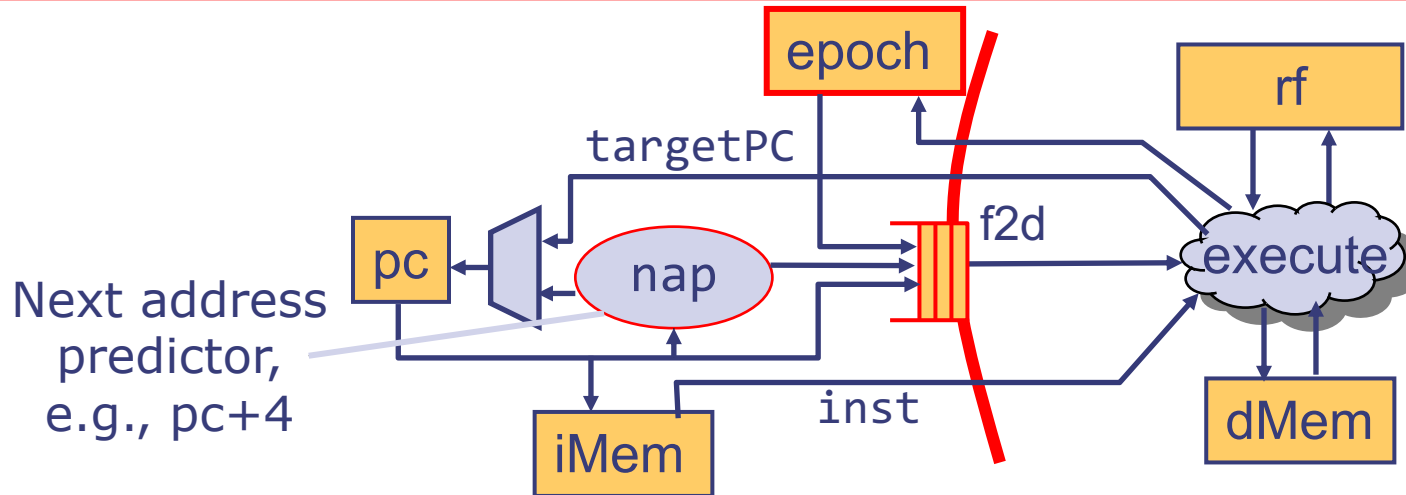
How do we detect a misprediction?

- We initiate a fetch for the instruction at pc, and make a prediction for the next pc (ppc)
- The instruction at pc carries the prediction (ppc) with it as it flows through the pipeline
- At the Execute stage we know the real next pc. It is a misprediction if the next pc \neq ppc

What does it mean to squash a partially executed instruction?


- The instruction should have no effect on the processor state
 - must not update register file or pc
 - must not launch a Store
- These conditions are easy to ensure in our two-stage processor because there is at most one instruction in the Ex/LW state

Epoch: a method to manage control hazards



- Add an *epoch* register to the processor state
- The Execute stage changes the *epoch* whenever the pc prediction is wrong and sets the pc to the correct value
- The Fetch stage associates the current *epoch* to every instruction sent to the Execute stage
- The epoch of the instruction is examined when it is ready to execute. If the processor epoch has changed the instruction is thrown away


Pipelining Decode and Execute



- Execute step probably has the longest propagation delay (decode + register-file read + execute)
- Separate Execute into two stages:
 - Decode and register-file-read
 - Execute – including the initiation of memory instructions
- This introduces a new problem known as a *Data Hazard*, that is, the register file, when it is read, may have stale values

Three stage pipeline


data hazard



I_0	R1	\leftarrow	Ld	R2
I_1	R4	\leftarrow	R1+R2	
I_2	...			

mutually
exclusive

Fetch
Decode
Ex
LW



- I_1 must be stalled until I_0 updates the register file, i.e., the data hazard disappears \Rightarrow need a mechanism to stall
- The data hazard will disappear as pipeline drains

Complication: the stalled instruction may
be a wrong-path instruction

next lecture

Data Hazard

- Data hazard arises when a source register of the fetched instruction matches the destination register of an instruction already in the pipeline
- Both the source and destination registers must be valid for a hazard to exist

Our current decoder does not record the validity of the source fields. However, the decoder can be changed easily to attach a valid field to each source register


Dealing with data hazards

(aka read-after-write (RAW) hazard)

- Introduce a *Scoreboard* -- a data structure to keep track of the destinations of the instructions in the pipeline beyond the Decode stage
 - Initially the scoreboard is empty
- Compare sources of an instruction when it is decoded with the destinations in the scoreboard
- Stall the Decode from dispatching the instruction to Execute if there is a RAW hazard
- When the instruction is dispatched, enter its destination in the scoreboard
- When an instruction completes, delete its source from the scoreboard


A stalled instruction will be unstalled when the RAW hazard disappears. This is guaranteed to happen as the pipeline drains.

Scoreboard



- *method insert(dst)*: inserts the destination of an instruction or Invalid in the scoreboard when the instruction is decoded
- *method search1(src)*: searches the scoreboard for a data hazard, i.e., a dst that matches src
- *method search2(src)*: same as *search1*
- *method remove*: deletes the oldest entry when an instruction commits

Two designs for scoreboard



versus

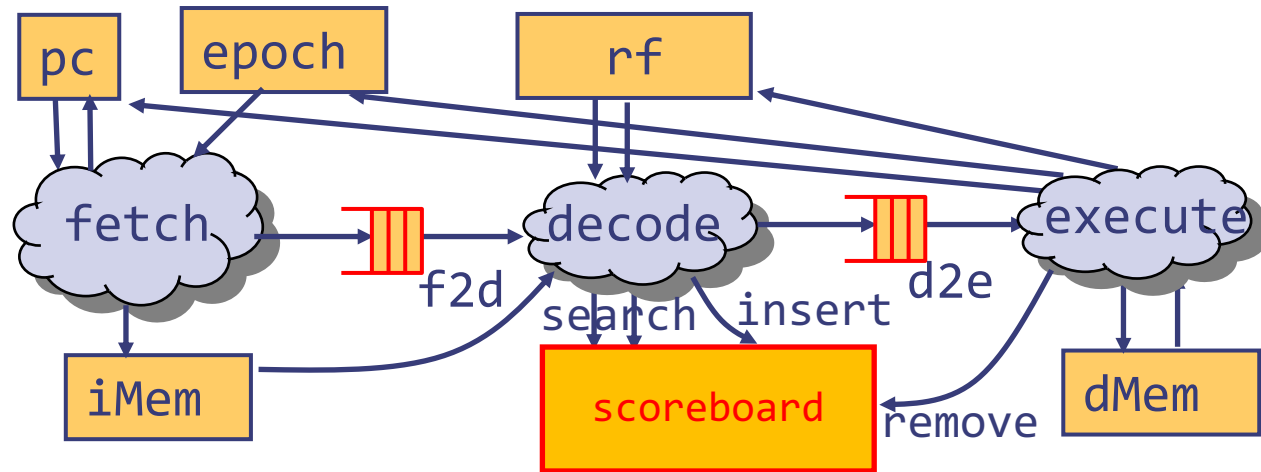


- A fifo of depth equal to the number of pipeline stages in Execute
- *Insert*: enq (dst)
- *Remove*: deq
- *Search*: compare source against each entry

- One Boolean flag for each register (Initially all False)
- *Insert*: set the flag for register rd to True (block if it is already True)
- *Remove*: set the flag for register rd to False
- *Search*: Return the value of the flag for the source register

Counter design takes less hardware, especially for deep pipelines, and is more efficient because it avoids searching each element of the fifo


Scoreboard in the pipeline



- If search by Decode does not see an instruction in scoreboard, then that instruction must have updated the state
- Thus, when an instruction is removed from the scoreboard, its updates to Register File must be visible to the subsequent register reads in Decode
 - remove and wr should happen simultaneously
 - search, and rd1, rd2 should happen simultaneously

This will require a bypass register file

Bypassing



- Bypassing is a technique to reduce the number of stalls (that is, the number of cycles) by providing extra data paths between the producer of a value and its consumer
- Bypassing introduces new combinational paths and this can increase combinational delay (and hence the clock period) and area
- The effectiveness of a bypass is determined by how often it is used

Normal vs Bypass Register File

```
module mkRFile(RFile);  
  Vector#(32,Reg#(Data)) rfile <- replicateM(mkReg(0));  
  
  method Action wr(RIndx rindx, Data data);  
    if(rindx!=0) rfile[rindx] <= data;  
  endmethod  
  
  method Data rd1(RIndx rindx) = rfile[rindx];  
  method Data rd2(RIndx rindx) = rfile[rindx];  
endmodule
```

$\{rd1, rd2\} < wr$

Can we design a bypass register file so that:

$wr < \{rd1, rd2\}$

Processor Performance

$$\frac{\text{Time}_{\text{Program}}}{\text{Program}} = \frac{\text{Instructions}_{\text{Program}}}{\text{Program}} \cdot \frac{\text{Cycles}_{\text{Instruction}}}{\text{Instruction}} \cdot \frac{\text{Time}_{\text{Cycle}}}{\text{Cycle}}$$

$\text{CPI} \quad t_{\text{clk}}$

- Pipelining lowers t_{clk} . What about CPI?
- $\text{CPI} = \text{CPI}_{\text{ideal}} + \text{CPI}_{\text{hazard}}$
 - $\text{CPI}_{\text{ideal}}$: cycles per instruction if no stall
- $\text{CPI}_{\text{hazard}}$ contributors
 - Data hazards: long operations, cache misses
 - Control hazards: branches, jumps, exceptions