


Module Interfaces and Concurrency

Arvind
Computer Science and Artificial Intelligence Laboratory
M.I.T.

Design Alternatives:

Latency and Throughput


Combinational
(C)




Suppose the latency of each

- f_i is 5ns
- FIFO is 1ns
- mux is 1ns
- register is 1ns

Pipeline
(P)



Folded:
Reuse a
block (F)




Area: $F < C < P$

- Latency: Total time to process 1 element
- Throughput: Rate at which new elements can be processed

| | C | P | F |
|-------------------|--------------------------|----------------------------|------------------------------------|
| Latency (ns) | $1 + 3 \cdot 5 + 1 = 17$ | $1 + 3 \cdot (5 + 1) = 19$ | $1 + 3 \cdot (1 + 5 + 1) + 1 = 23$ |
| Throughput (1/ns) | $1/17$ | $1/(5 + 1) = 1/6$ | $1/23$ |
| Clock period | 16 | 6 | 7 |

A pipelined system



Without concurrent execution it is hardly a pipelined system

```
rule stage1;  
  fifo1.enq(f0(inQ.first));  
  inQ.deq;  endrule  
rule stage2;  
  fifo2.enq(f1(fifo1.first));  
  fifo1.deq; endrule  
rule stage3;  
  outQ.enq(f2(fifo2.first));  
  fifo2.deq; endrule
```

- When can stage1 rule fire?
 - *inQ has an element*
 - *fifo1 has space*
- Can these 3 rules execute concurrently?
 - *Yes, but it must be possible to do enq and deq in a fifo simultaneously*
- Can stage1 and stage3 execute concurrently?
 - *Yes, even if enq and deq cannot be done simultaneously in a fifo*

Multirule Systems

- Most systems we have seen so far had multiple rules but only one rule was ready to execute at any given time (pair-wise mutually exclusive rules)
- Consider a system where multiple rules can be ready to execute at a given time
 - When can two such rules be executed together?
 - What does the synthesized hardware look like for concurrent execution of rules?

One-rule-at-a-time semantics of Bluespec

Repeatedly:

- ◆ Select any rule that is ready to execute
- ◆ Compute the state updates
- ◆ Make the state updates

Any legal behavior of a Bluespec program can be explained by observing the state updates obtained by applying one rule at a time

However, for performance we execute multiple rules concurrently whenever possible

Concurrent execution of rules

- Two rules can execute concurrently, if concurrent execution would not cause a double-write error, *and*
- The final state can be obtained by executing rules one-at-a-time in some sequential order
- Can these rules execute concurrently without violating the one-rule-at-a-time-semantics?

Example 1

```
rule ra;  
  x <= x+1;  
endrule  
rule rb;  
  y <= y+2;  
endrule
```

Example 2

```
rule ra;  
  x <= y+1;  
endrule  
rule rb;  
  y <= x+2;  
endrule
```

Example 3

```
rule ra;  
  x <= y+1;  
endrule  
rule rb;  
  y <= y+2;  
endrule
```

Concurrent Execution

Example 1

```
rule ra;  
   $x^{t+1} \leq x^{t+1};$   
endrule  
rule rb;  
   $y^{t+1} \leq y^{t+2};$   
endrule
```

Example 2

```
rule ra;  
   $x^{t+1} \leq y^{t+1};$   
endrule  
rule rb;  
   $y^{t+1} \leq x^{t+2};$   
endrule
```

Example 3

```
rule ra;  
   $x^{t+1} \leq y^{t+1};$   
endrule  
rule rb;  
   $y^{t+1} \leq y^{t+2};$   
endrule
```

Final value of (x,y) given the initial values (0,0)

Concurrent
Execution

Ex 1
(1,2)

Ex 2
(1,2)

Ex 3
(1,2)

ra < rb

rb < ra

Executing ra before rb ($ra < rb$)

Example 1

```
rule ra;  
   $x^{t+1} \leq x^{t+1}$ ;  
endrule  
rule rb;  
   $y^{t+2} \leq y^{t+1}+2$ ;  
endrule
```

Example 2

```
rule ra;  
   $x^{t+1} \leq y^{t+1}$ ;  
endrule  
rule rb;  
   $y^{t+2} \leq x^{t+1}+2$ ;  
endrule
```

Example 3

```
rule ra;  
   $x^{t+1} \leq y^{t+1}$ ;  
endrule  
rule rb;  
   $y^{t+2} \leq y^{t+1}+2$ ;  
endrule
```

Final value of (x,y) given the initial values $(0,0)$

| | Ex 1 | Ex 2 | Ex 3 |
|----------------------|---------|---------|---------|
| Concurrent Execution | $(1,2)$ | $(1,2)$ | $(1,2)$ |
| $ra < rb$ | $(1,2)$ | $(1,3)$ | $(1,2)$ |
| $rb < ra$ | | | |

For any x , if there is no x^{t+1} defined, then $x^{t+1} = x^t$

Executing rb before ra ($rb < ra$)

Example 1

```
rule ra;  
   $x^{t+2} \leq x^{t+1}+1$ ;  
endrule  
rule rb;  
   $y^{t+1} \leq y^t+2$ ;  
endrule
```

Example 2

```
rule ra;  
   $x^{t+2} \leq y^{t+1}+1$ ;  
endrule  
rule rb;  
   $y^{t+1} \leq x^t+2$ ;  
endrule
```

Example 3

```
rule ra;  
   $x^{t+2} \leq y^{t+1}+1$ ;  
endrule  
rule rb;  
   $y^{t+1} \leq y^t+2$ ;  
endrule
```

Final value of (x,y) given the initial values $(0,0)$

| | Ex 1 | Ex 2 | Ex 3 |
|----------------------|-------|-------|-------|
| Concurrent Execution | (1,2) | (1,2) | (1,2) |
| $ra < rb$ | (1,2) | (1,3) | (1,2) |
| $rb < ra$ | (1,2) | (3,2) | (3,2) |

Can these rules execute concurrently?

(without violating the one-rule-at-a-time-semantics)

Example 1

```
rule ra;  
  x <= x+1;  
endrule  
rule rb;  
  y <= y+2;  
endrule
```

Example 2

```
rule ra;  
  x <= y+1;  
endrule  
rule rb;  
  y <= x+2;  
endrule
```

Example 3




```
rule ra;  
  x <= y+1;  
endrule  
rule rb;  
  y <= y+2;  
endrule
```

Final value of (x,y) given the initial values (0,0)

Concurrent
Execution

ra < rb

rb < ra



Conflict Matrix (CM)

BSV compiler generates the pairwise conflict information

Example 1

```
rule ra;  
  x <= x+1;  
endrule  
rule rb;  
  y <= y+2;  
endrule
```

| | ra | rb |
|----|----|----|
| ra | C | CF |
| rb | CF | C |

Example 2

```
rule ra;  
  x <= y+1;  
endrule  
rule rb;  
  y <= x+2;  
endrule
```

| | ra | rb |
|----|----|----|
| ra | C | C |
| rb | C | C |

Example 3

```
rule ra;  
  x <= y+1;  
endrule  
rule rb;  
  y <= y+2;  
endrule
```

| | ra | rb |
|----|----|----|
| ra | C | < |
| rb | > | C |

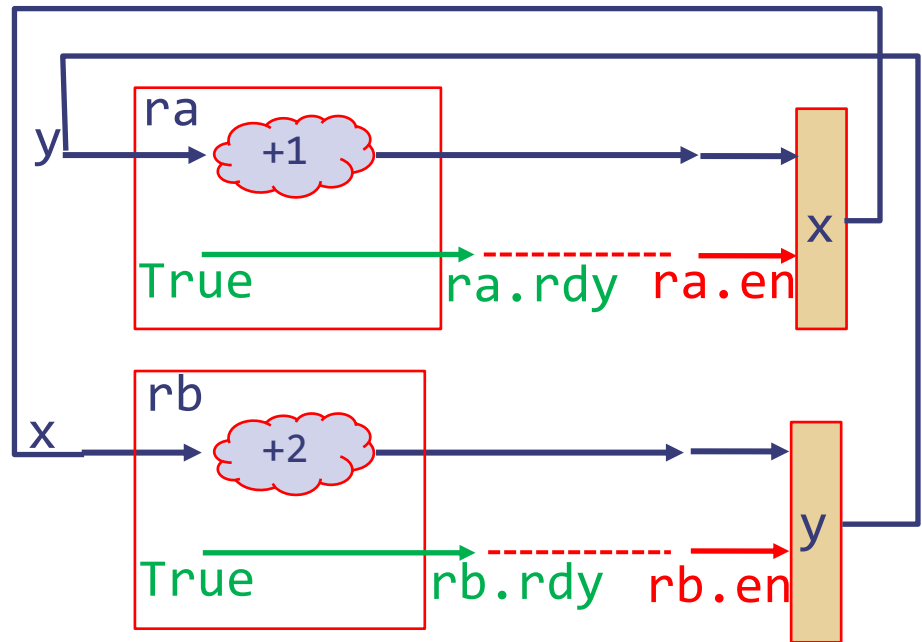
- C : rules can't be executed concurrently
- CF: rules can be executed concurrently; the net effect is the same as if ra executed before rb (ra<rb) or (rb<ra)
- ra < rb : rules can be executed concurrently; the net effect is as if ra executed before rb

Using *conflict* (CM) information in hardware synthesis

Concurrent rule execution

Example 2

```
rule ra;  
  x <= y+1;  
endrule  
rule rb;  
  y <= x+2;  
endrule
```



- Suppose we connect **ra.rdy** to **ra.en**, and **rb.rdy** to **rb.en**
- This circuit will execute rules ra and rb concurrently
- This circuit is correct only if rules ra and rb do not conflict
- But in this example rules ra and rb do conflict!

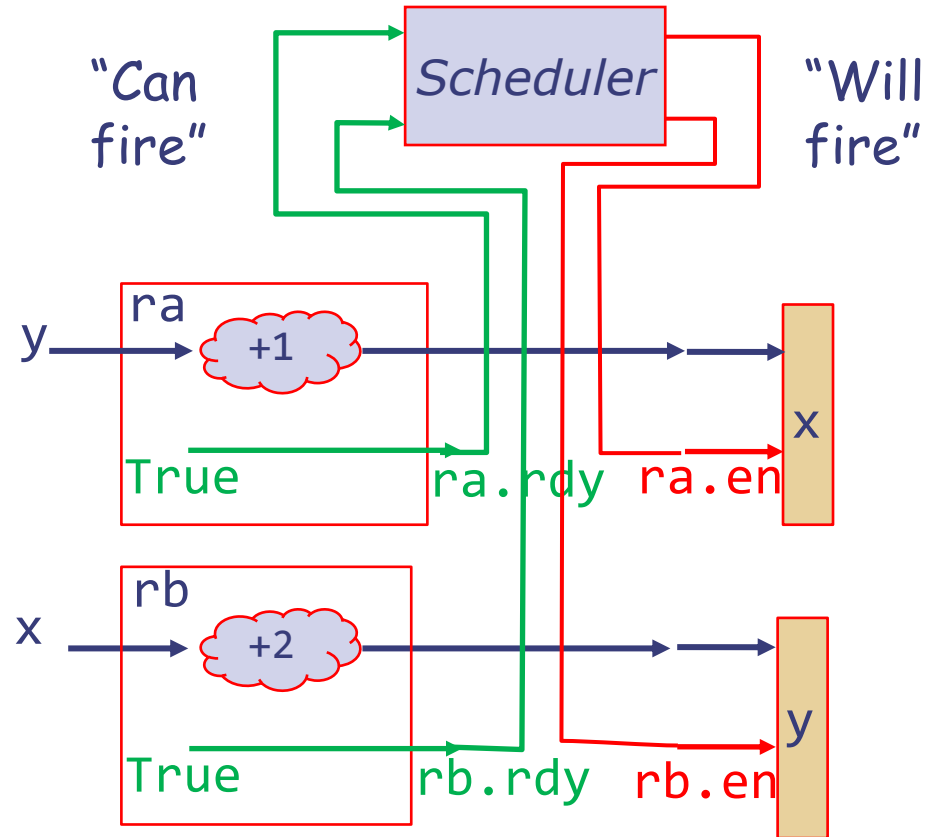


Need for a scheduler

Example 2

```
rule ra;  
  x <= y+1;  
endrule  
rule rb;  
  y <= x+2;  
endrule
```


- Guards (**rdy** signals) of all rules are fed to a scheduler
- Using the CM, the scheduler lets only non-conflicting rules proceed
 - Scheduler is a pure combinational circuit with a small number of gates
 - A correct but low performance scheduler may schedule only one rule at a time



Example schedulers


Example1

| | ra | rb |
|----|----|----|
| ra | C | CF |
| rb | CF | C |




Example2

| | ra | rb |
|----|----|----|
| ra | C | C |
| rb | C | C |



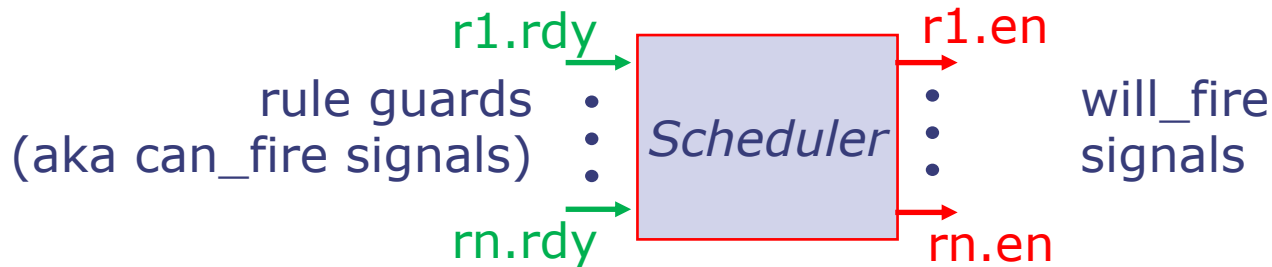
Example3

| | ra | rb |
|----|----|----|
| ra | C | < |
| rb | > | C |



The effect $ra < rb$ would be achieved automatically

The rule scheduler



- Guards ($r1.rdy \dots rn.rdy$) of many rules may be true simultaneously, and some of them may conflict
- BSV compiler constructs a combinational scheduler circuit with the following property:

for all pairs of rules i and j , if $r^i.en$ and $r^j.en$ are true then the corresponding $r^i.rdy$ and $r^j.rdy$ must be true and rules i and j *must not conflict* with each other

Takeaway


- One-rule-at-a-time semantics are important to understand the legal behaviors of a system
- Efficient hardware for multi-rule system requires that many rules execute in parallel without violating the one-rule-at-a-time semantics
- BSV compiler builds a scheduler circuit to execute as many rules as possible concurrently
- For high-performance designs we have to worry about the CM characteristics of our modules

More on this topic later in the course

We are not done yet!

```
rule stage1;
  fifo1.enq(f0(inQ.first));
  inQ.deq;   endrule
rule stage2;
  fifo2.enq(f1(fifo1.first));
  fifo1.deq; endrule
rule stage3;
  outQ.enq(f2(fifo2.first));
  fifo2.deq; endrule
```

- These rules can execute concurrently, only if fifos allow concurrent enq and deq
- In our one-element fifo design, enq and deq are mutually exclusive!
- We will design better FIFO's later




No
pipelining

| | enq | deq | first |
|-------|-----|-----|-------|
| enq | C | ME | ME |
| deq | ME | C | > |
| first | ME | < | CF |

| | enq | deq | first |
|-------|-----|-----|-------|
| enq | C | CF | CF |
| deq | CF | C | > |
| first | CF | < | CF |

Take-home problem

- Draw a hardware circuit for this design, ignoring the internals of the fifo design



Hint

- Draw the guard for each rule
- Assume an enable signal for each rule
- Connect the rule ready and enable signals to a scheduler

```
rule stage1;  
    fifo.enq(f0(inQ.first));  
    inQ.deq;  
endrule  
rule stage2;  
    outQ.enq(f1(fifo.first));  
    fifo.deq;  
endrule
```