

# Operating Systems

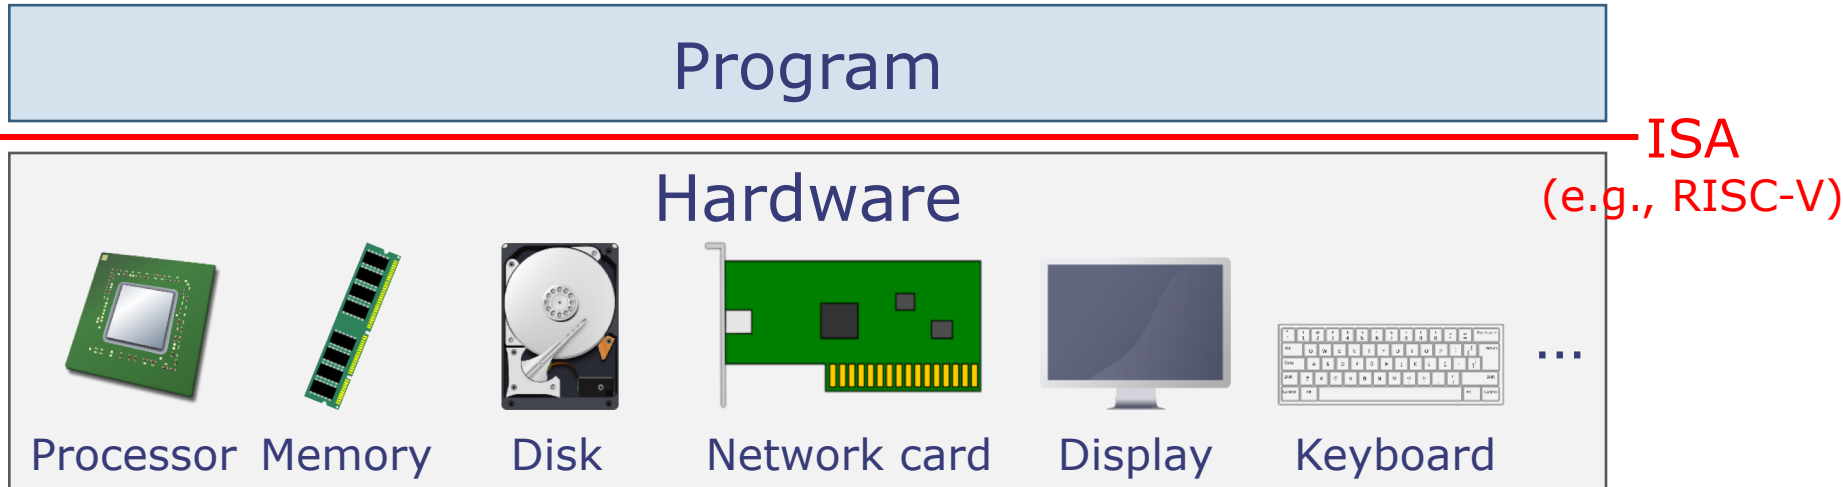
*Arvind*

Computer Science & Artificial Intelligence Lab  
M.I.T.

Quiz 2 tonight 7:30-9:30pm  
Room 50-340

# 6.004 So Far: Single-User Machines

---

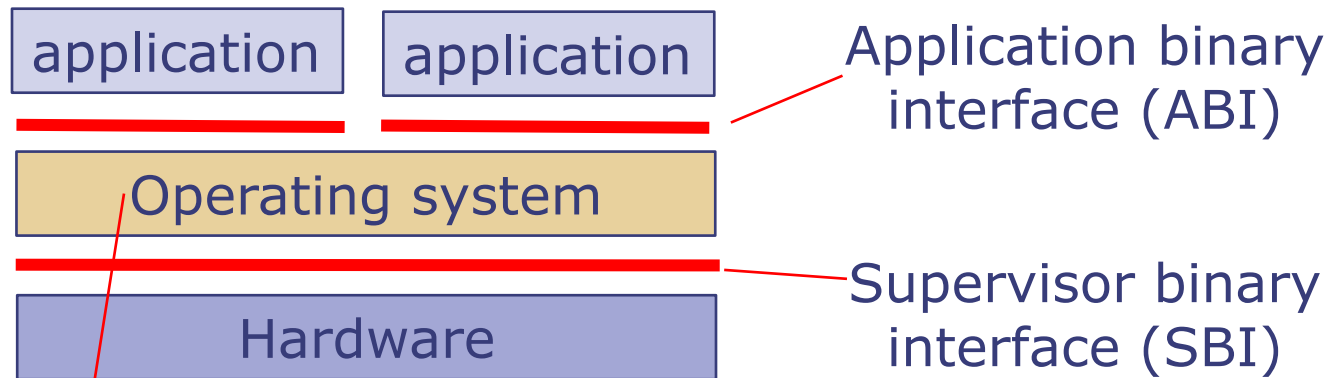


- Hardware executes a single program
- This program has direct and complete access to all hardware resources in the machine
- The instruction set architecture (ISA) is the interface between software and hardware

This picture is too simplistic.

# Operating Systems

---



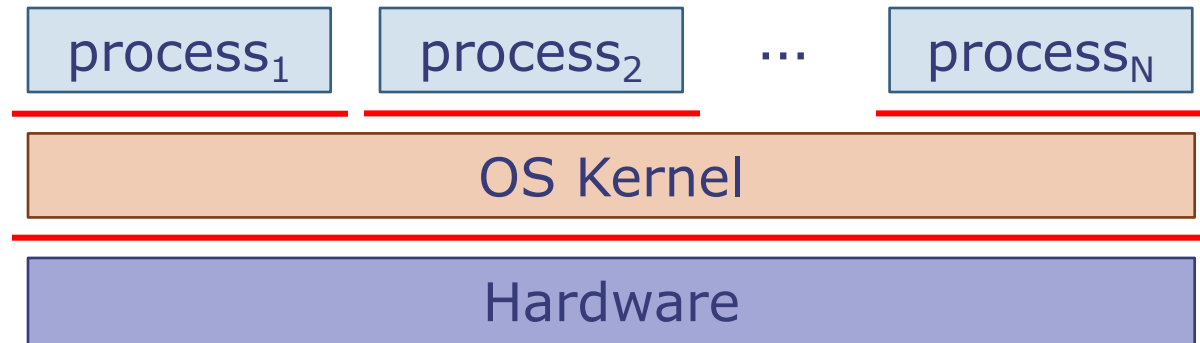
- Modern computers cannot function without some resident programs (“services”), which are shared by all users
  - Services: programs for accessing I/O devices like, keyboard, mouse, display, disks, printers, network, ...
  - I/O devices have very long access latency as compared to registers and memories (milliseconds as opposed to microseconds)
  - It is common to employ “multiprogramming” and switch to another program while waiting for an I/O operation to complete

**Supporting Operating Systems creates a host of new problems in processor design**

Nomenclature:

# Process, Program, Processor

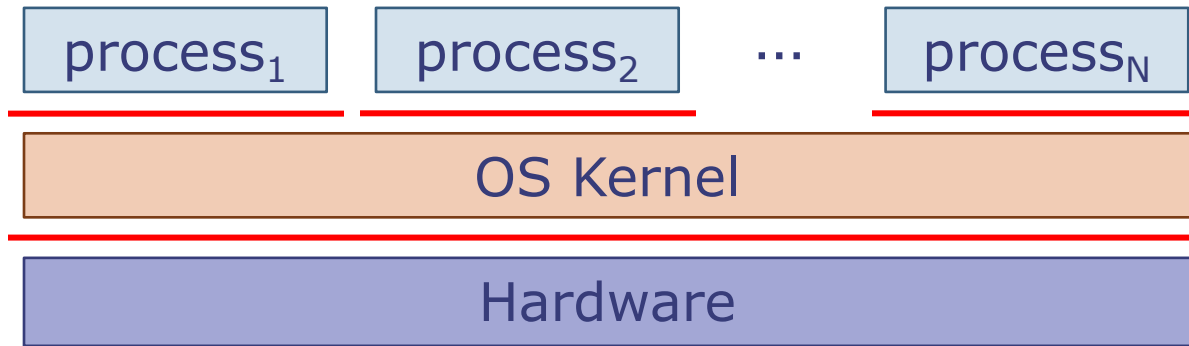
---



- A program is a collection of instructions and data
- A processor is the hardware to execute programs
- A **process** is an instance of a program that is being executed
  - Includes program, processor state (registers, memory), and I/O device state
- The **OS Kernel** is a special process with extra privileges


# Goals of Operating Systems

---




- **Protection** and privacy: Processes cannot access each other's data
- **Abstraction**: OS hides details of underlying hardware
  - e.g., processes open and access files instead of issuing raw commands to the disk
- **Resource management**: OS controls, i.e., allocates and schedules, how processes share hardware (CPU, memory, disk, etc.)

# Private Address Space per Process




- Each process is allocated space in physical memory by the OS kernel
- A process is not allowed to access the memory of other processes or the OS
- It is convenient if the program and data addresses don't depend upon where the process is allocated in the memory
  - requires dynamic address translation

next lecture




# Scheduling

- The OS kernel **schedules processes** into the CPU
  - Each process is given a fraction of CPU time
  - A process cannot use more CPU time than allowed



# System Calls

- The OS kernel lets processes invoke system services (e.g., access files or network sockets) via **system calls** (special instruction)
  - Key Board input
  - Mouse
  - Display
  - File Access
- ABI defines how process and kernel pass arguments and results; typically, similar conventions as a function call:
  - **System call number** in a7
  - Other arguments in a0-a6
  - Results in a0-a1 (or in memory)
  - All registers are preserved

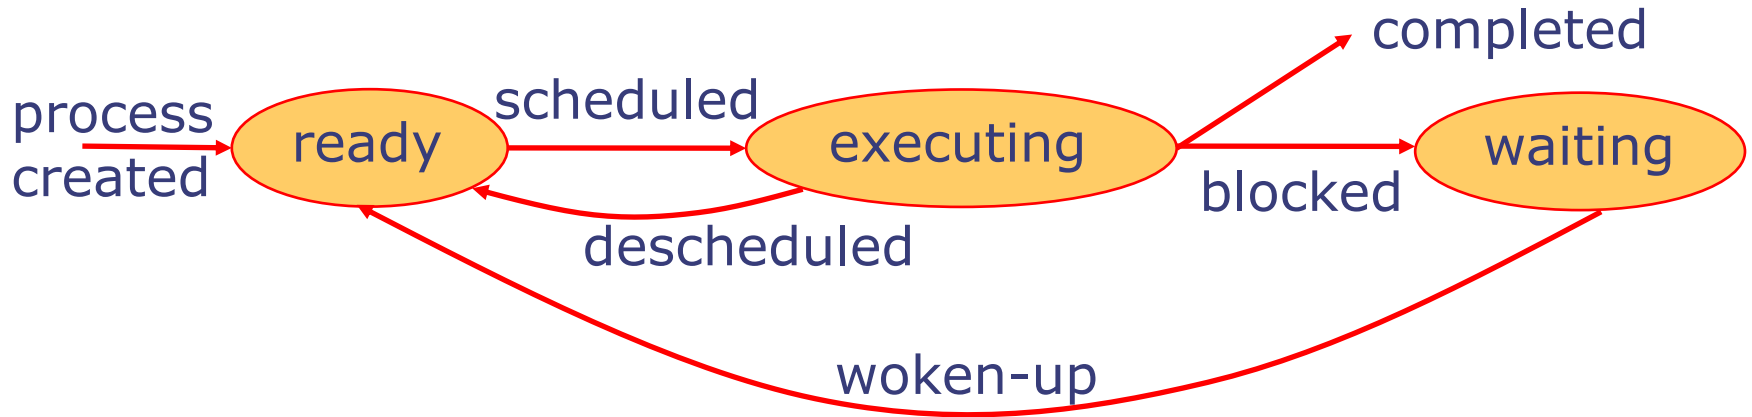


Lab 8



# Process Life Cycle

---



- OS maintains a list of all processes and their status {ready, executing, waiting}
  - A process is scheduled to run for a specified amount of CPU time or until completion
  - If a process invokes a system call that cannot be satisfied immediately (e.g., a file read that needs to access disk), it is *blocked* and put in the *waiting* state
  - When the waiting condition has been satisfied, the waiting process is woken up and put in the ready list

# ISA Extensions to Support OS

---


- Two modes of execution: **user** and **supervisor**
  - OS kernel runs in supervisor mode
  - All other processes run in user mode
- **Privileged instructions and registers** that are only available in supervisor mode
- **Interrupts and exceptions** to safely transition from user to supervisor mode Today
- **Virtual memory** to provide private address spaces and abstract the storage resources of the machine Next lecture

These ISA extensions work only if hardware and software (OS) agree on a common set of conventions!

# Exceptions

---

- Exception: Event that needs to be processed by the OS kernel. The event is usually unexpected or rare.



# Causes for Exceptions

---

- The terms exception and interrupt are often used interchangeably, with a minor distinction:
- **Exceptions** usually refer to **synchronous events**, generated by the process itself (e.g., illegal instruction, divide-by-0, illegal memory address, system call)
- **Interrupts** usually refer to **asynchronous events**, generated by I/O devices (e.g., timer expired, keystroke, packet received, disk transfer complete)
- We use exception to encompass both types of events, and use synchronous exception for synchronous events

# Handling Exceptions

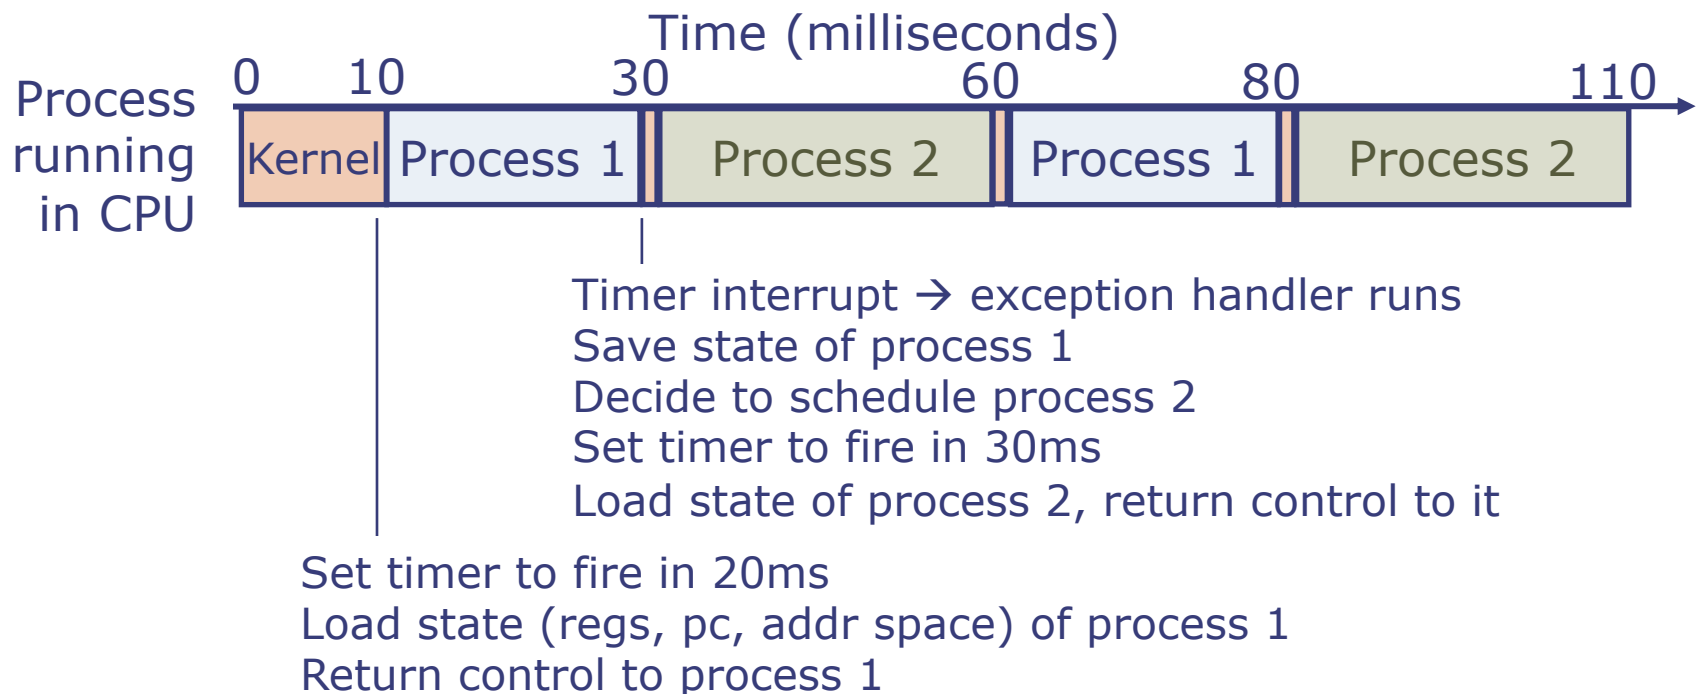
---

- When an exception happens, the processor:
  - Stops the current process at instruction  $I_i$ , completing all the instructions up to  $I_{i-1}$  (*precise exceptions*)
  - Saves the PC of instruction  $I_i$  and the reason for the exception in special (privileged) registers
  - Enables supervisor mode, disables interrupts, and transfers control to a pre-specified exception handler PC
- After the OS kernel handles the exception, it returns control to the process at instruction  $I_i$ 
  - Exception is transparent to the process!
- If the exception is due to an illegal operation by the program that cannot be fixed (e.g., an illegal memory access), the OS aborts the process

# Case Study 1: CPU Scheduling

Enabled by timer interrupts

- The OS kernel **schedules processes** into the CPU
  - Each process is given a fraction of CPU time
  - A process cannot use more CPU time than allowed
- Key enabling technology: Timer interrupts
  - Kernel sets timer, which raises an interrupt after a specified time



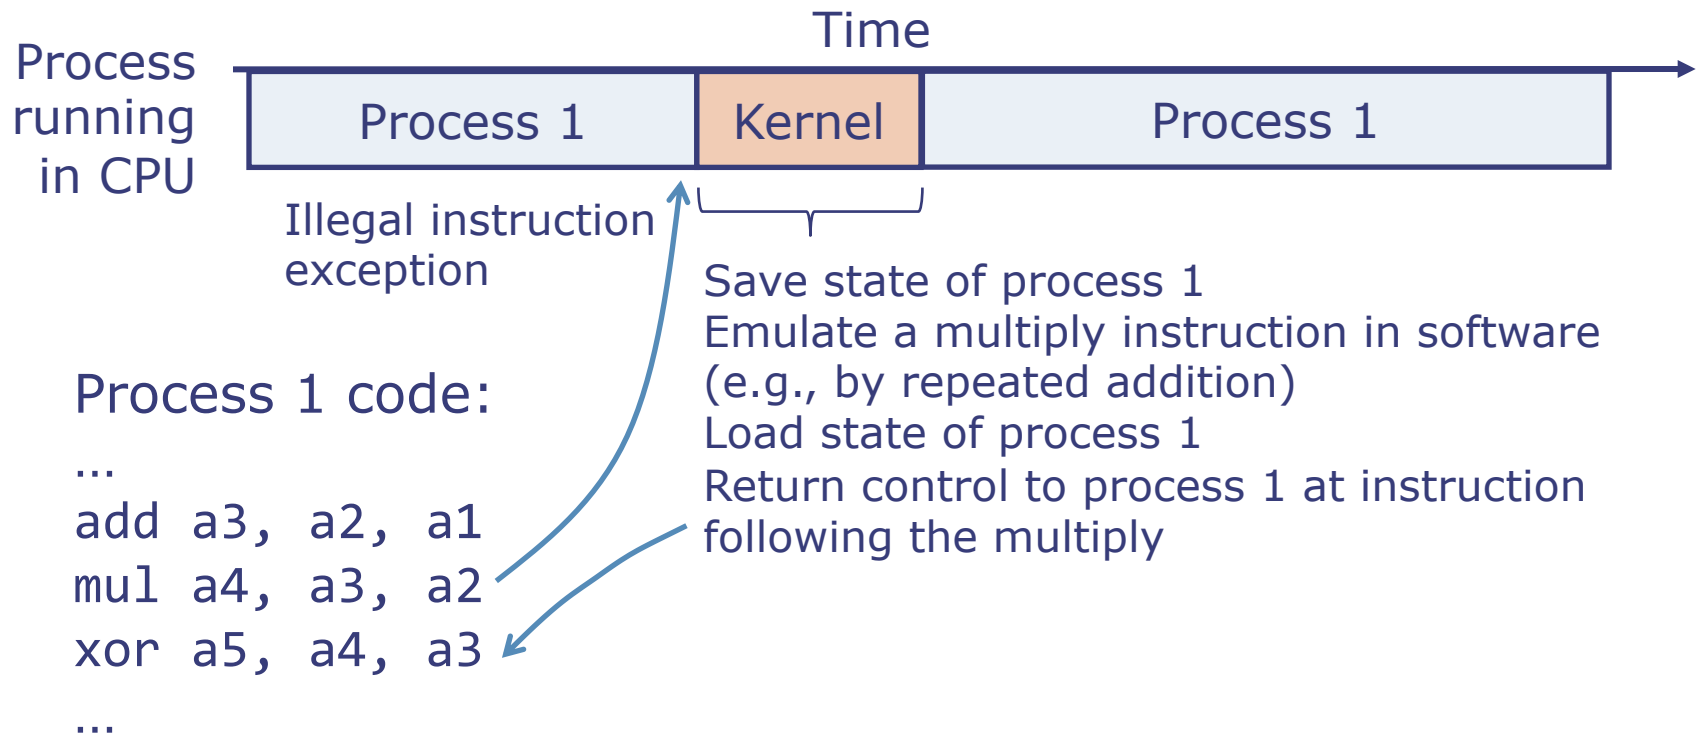
# Case Study 2: Emulating Instructions

Enabled by illegal instruction exceptions

---

- `mul x1, x2, x3` is an instruction in the RISC-V 'M' extension ( $x1 := x2 * x3$ )
  - If 'M' is not implemented, this is an illegal instruction
- What happens if we run code from an RV32IM machine on an RV32I machine?
  - `mul` causes an illegal instruction exception
- The exception handler can take over and abort the process... but it can also emulate the instruction!

# Emulating Unsupported Instructions



- Result: Program believes it is executing in a RV32IM processor, when it's actually running in a RV32I
  - *Any problem?* **Much slower than a hardware multiply**



# RISC-V Exception Handling


---

- RISC-V provides several privileged registers, called control and status registers (CSRs), e.g.,
  - **mepc**: exception PC
  - **mcause**: cause of the exception (interrupt, illegal instr, ...)
  - **mtvec**: address of the exception handler
  - **mstatus**: status bits (privilege mode, interrupts enabled, ...)
  - ...
- RISC-V also provides privileged instructions, e.g.,
  - **csrr** and **csrw** to read/write CSRs
  - **mret** to return from the exception handler to the process

Executing a privileged instruction in user mode causes an exception; a user process cannot take over the system

# Typical Exception Handler Structure

- A small common handler (CH) written in assembly + many exception handlers (EHs), one for each cause (typically written in normal C code)
- Common handler is invoked on every exception:
  1. Saves registers x1-x31, mepc into known memory locations
  2. Passes mcause, process state to the right EH to handle the specific exception/interrupt
  3. EH returns which process should run next (could be the same or a different one)
  4. CH loads x1-x31, mepc from memory for the right process
  5. CH executes mret, which sets pc to mepc, disables supervisor mode, and re-enables interrupts



# Common Exception Handler

## RISC-V Assembly code

---

```
common_handler: // entry point for exception handler
// save x1 to mscratch to free up a register
csrw mscratch, x1 // write x1 to mscratch CSR
// get the pointer for current process's state
lw x1, curProcState
// save registers x2-x31
sw x2, 8(x1)
sw x3, 12(x1)
...
sw x31, 124(x1)
// now registers x2-x31 are free for the kernel
// save original x1 (now in mscratch)
csrr t0, mscratch
sw t0, 4(x1)
// finally, save mepc
csrr t1, mepc
sw t1, 0(x1)
```

# Common Exception Handler *cont.*

## Calling EH\_Dispatcher and returning

---

```
common_handler:
    ... // we have saved the state of the process
    // pass interrupted process state and cause to eh_dispatcher
    mv a0, x1 // arg 0: pointer interrupted process state
    csrr a1, mcause // arg 1: exception cause
    lw sp, kernelSp // use the kernel's stack
    // calls the appropriate handler
    jal eh_dispatcher
    // returns address of state of process to schedule
    // restore return PC in mepc
    lw t0, 0(a0)
    csrwr mepc, t0
    // restore x1-x31
    mv x1, a0
    lw x2, 8(x1); lw x3, 12(x1); ...; lw x31, 124(x1)
    lw x1, 4(x1) // restore x1 last
    mret // return control to program
```

# EH Dispatcher (in C)

Dispatches to a specific handler based on “cause”

---


```
typedef struct {
    int pc;
    int regs[31]; // reg 0 is not passed
    ...
} ProcState;

ProcState* eh_dispatcher(ProcState* curProc, int cause) {
    if(cause == 0x02) // illegal instruction
        return illegal_eh(curProc, cause);
    else if(cause == 0x08)
        // system call, e.g, OS service “write” to file
        return syscall_ih(curProc, cause);
    else if (cause < 0) // external interrupt
        ...
}
```

# Summary

---

- Operating System goals:
  - Protection and privacy: Processes cannot access each other's data
  - Abstraction: OS hides details of underlying hardware
    - e.g., processes open and access files instead of issuing raw commands to disk
  - Resource management: OS controls how processes share hardware resources (CPU, memory, disk, etc.)
- Key enabling technologies:
  - User mode + supervisor mode w/ privileged instructions
  - Exceptions to safely transition into supervisor mode
  - Virtual memory to provide private address spaces and abstract the machine's storage resources



next lecture