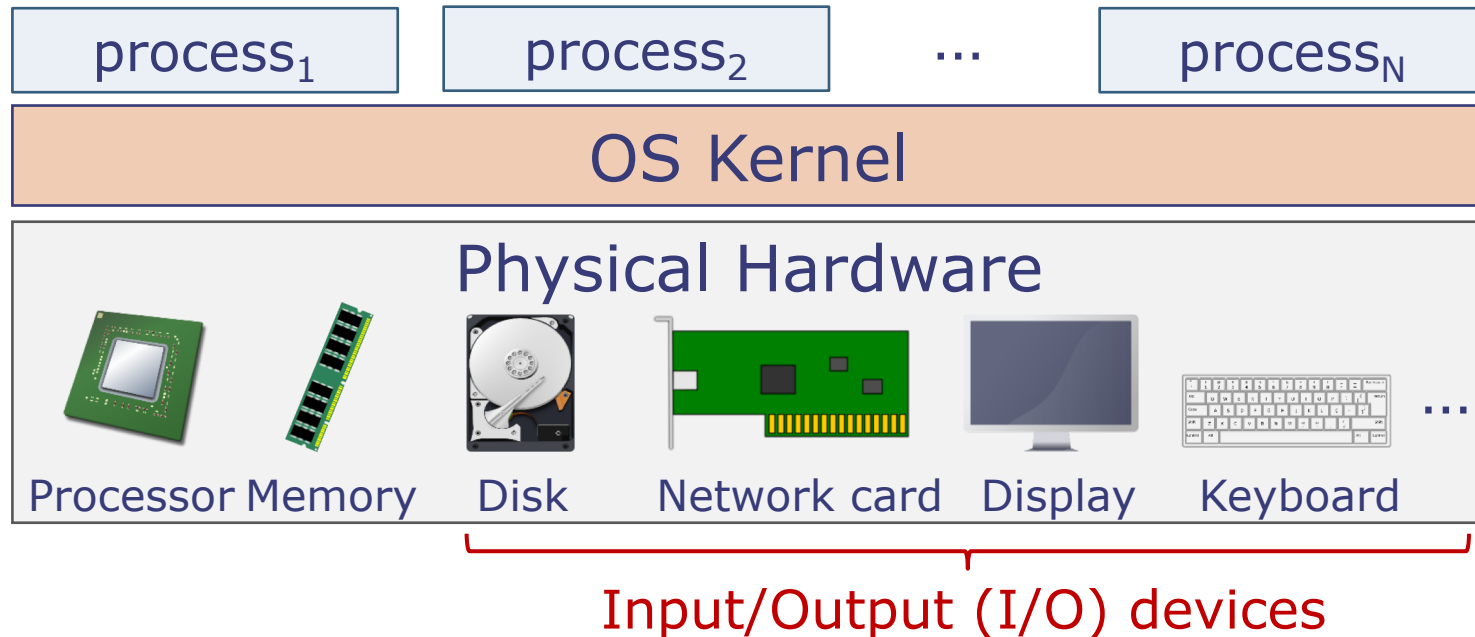




Communicating with I/O devices



- The system has shared **I/O registers** that both the processor and the I/O devices can read and write
- Key questions:
 - How to access I/O registers?
 - How do processor and device coordinate on each transfer?

Accessing I/O registers

- Option 1: Use special instructions
 - e.g., in and out instructions in x86
 - Inflexible, adds instructions → used rarely
- Option 2: Memory-mapped I/O (MMIO)
 - I/O registers are mapped to physical memory locations
 - Processor accesses them with loads and stores
 - These loads and stores should not be cached!



Coordinating I/O Transfers


- Option 1: Polling (synchronous)
 - Processor periodically reads the register associated with a specific I/O device
- Option 2: Interrupts (asynchronous)
 - Processor initiates a request, then moves to other work
 - When the request is serviced, the I/O device interrupts the processor
- *Pros of each approach?*

Polling is simple

Interrupts let the processor do useful computation while request is serviced

Example 1: Polling-based I/O

- Consider a simple character-based display
- Uses one I/O register with the following format:




- The **ready** bit (bit 31) is set to 1 only when the display is ready to print a character
- When the processor wants to display an 8-bit character, it writes it to **char** (bits 0-7) and sets the **ready** bit to 0
- After the display has processed the character, it sets the **ready** bit to 1

Let's see a demo!

Example 2: Interrupt-based I/O

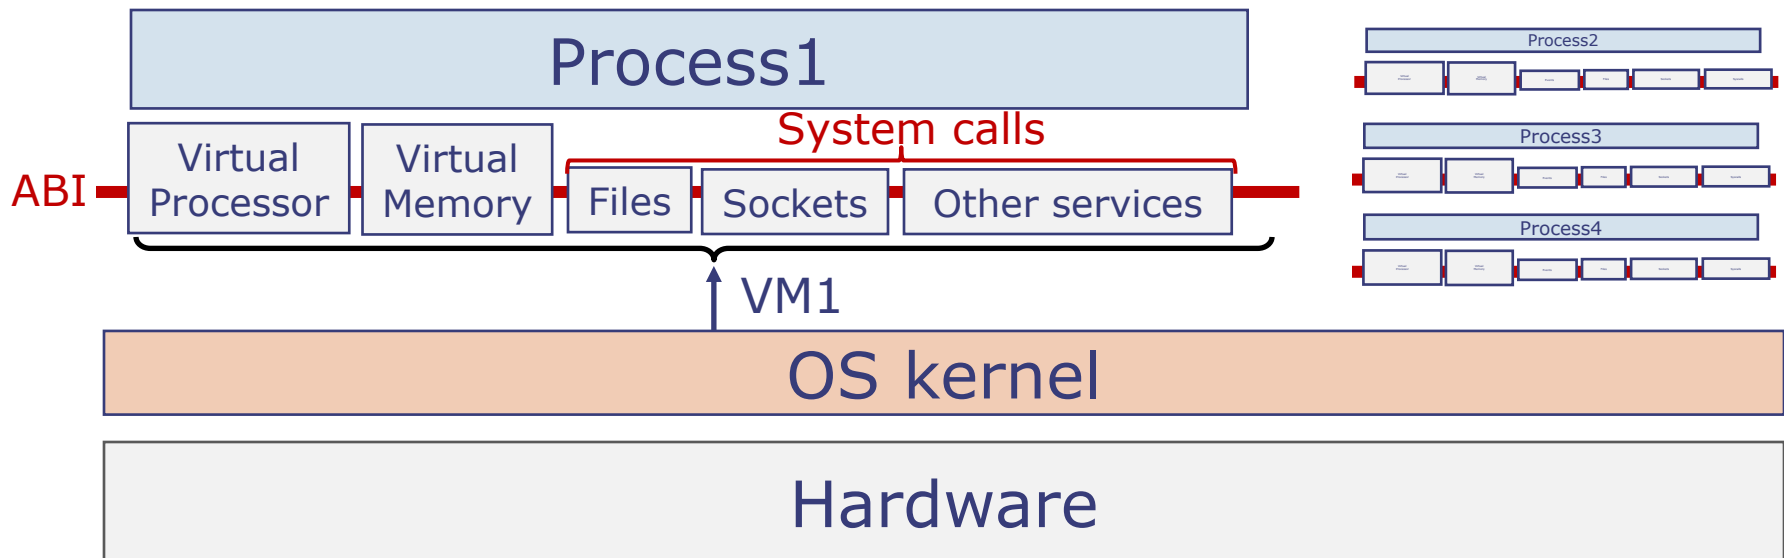
- Consider a simple keyboard that uses a single I/O register with a similar format:



- On a keystroke, the keyboard writes the typed character in `char`, sets `full` bit to 1, and **raises a keyboard interrupt**
- Interrupt handler reads `char` and sets `full` bit to 0, so the keyboard can deliver future keystrokes
- This works fine because keyboard is very slow compared to CPU. Faster devices use more sophisticated mechanisms (e.g., network cards write packets to main memory and use I/O registers only to indicate status of transfer)

Communicating with the OS

- The OS kernel lets processes invoke system services (e.g., access files) via **system calls**



- Processes invoke system calls by executing an instruction that causes an exception
 - Same mechanism as before!

System Calls in RISC-V

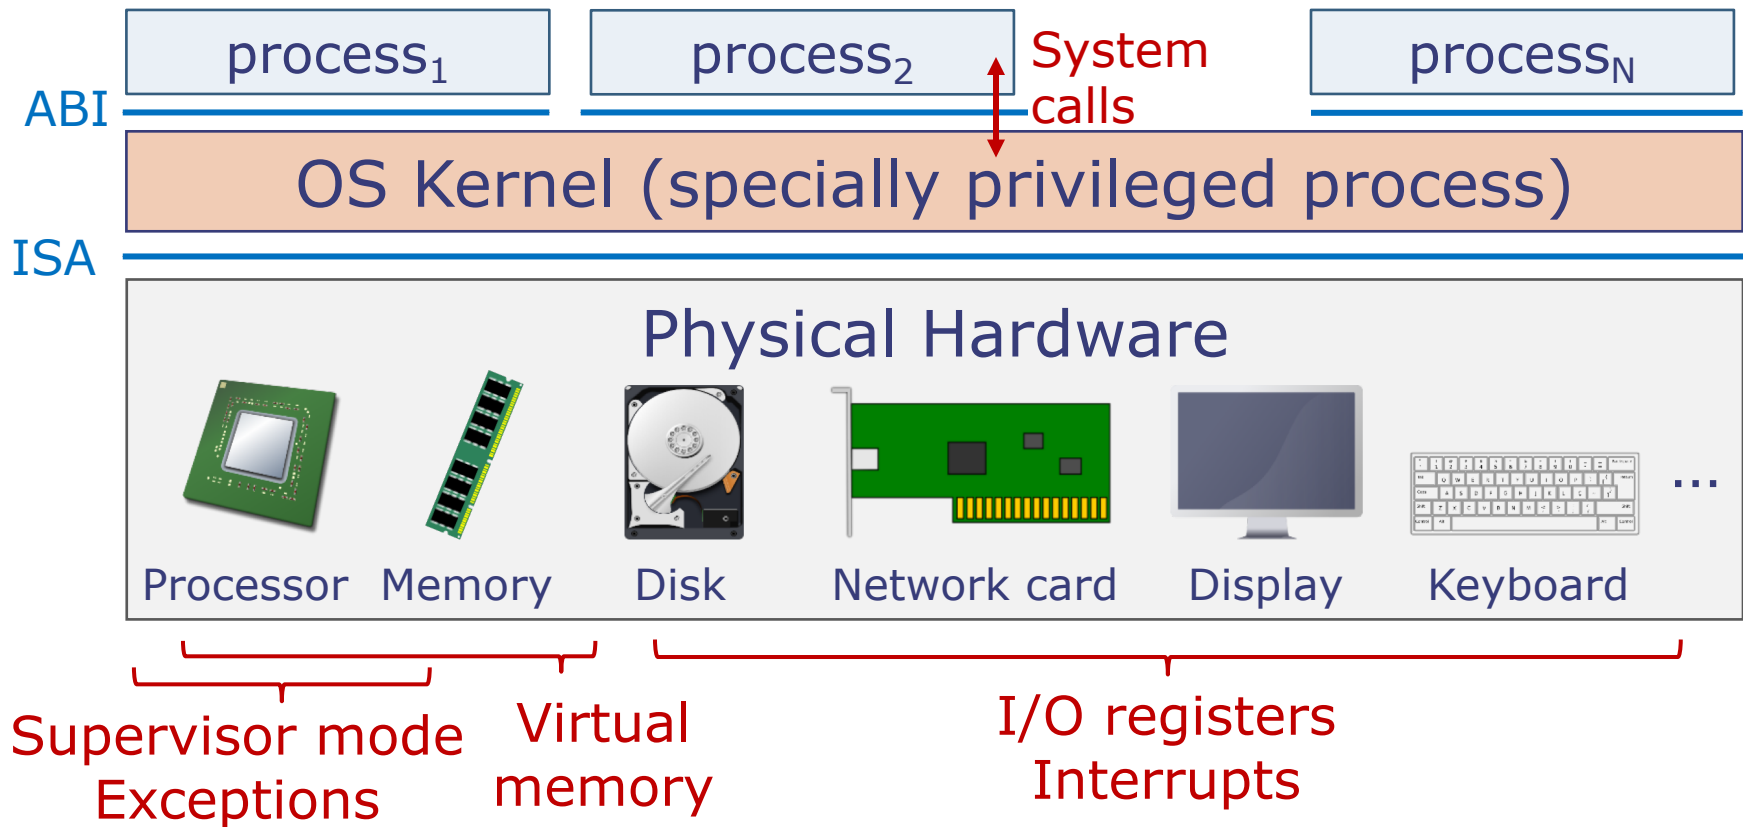
- `ecall` instruction causes an exception, sets `mcause` CSR to a particular value
- ABI defines how process and kernel pass arguments and results
- Typically, similar conventions as a function call:
 - **System call number** in `a7`
 - Other arguments in `a0-a6`
 - Results in `a0-a1` (or in memory)
 - All registers are preserved (treated as callee-saved)

Let's see a demo!

Typical System Calls

- Accessing files (`sys_open/close/read/write/...`)
 - Using network connections (`sys_bind/listen/accept/...`)
 - Managing memory (`sys_mmap/munmap/mprotect/...`)
 - Getting information about the system or process (`sys_gettime/getpid/getuid/...`)
 - Waiting for a certain event (`sys_wait/sleep/yield...`)
 - Creating and interrupting other processes (`sys_fork/exec/kill/...`)
 - ... and many more!
-
- Programs rarely invoke system calls directly. Instead, they are used by library/language routines
 - Some of these system calls may block the process!

Summary



- Want to learn more? Check out 6.033 or 6.828!