


Introduction to Pipelining

Arvind

Computer Science & Artificial Intelligence Lab
M.I.T.

Pipelining: A technique to increase the throughput of a system




- All stages of a pipeline run at the same clock speed
- *Throughput*: Rate at which items are processed
- *Latency*: Total time to process 1 item
- The notion of an item and associated answer depends on the application, e.g.,
 - Arithmetic pipeline for function f : input is number x , output is $f(x)$
 - Processor: item is an instruction
- The concept of pipelining is applicable to both combinational and sequential circuits, i.e., each f_i could be a combinational or sequential circuit

we will first study
combinational pipelines


Example:

Latency and Throughput




	Non pipelined	Pipelined
Clock period (ns)	$5+6+5=16$	6
Throughput (1/ns)	$1/16$	$1/6$
Latency (ns)	16	$3*6 = 18$


Pipelining combinational circuits



- Functionality of a combinational circuit can be specified as a pure function of inputs
- Correct pipelining must preserve the functionality



Incorrect Pipelining



Incorrect pipelining

	t_0	t_1	t_2	t_3
f	x_0	x_1		
g	x_0	x_1		
h	$(?, x_0)$	(x_0, x_1)		

$$\text{output} = h(f(x_i), g(x_{i+1}))$$

Every input to output path must have the same number of pipeline registers otherwise the computation would not divide into stages properly!

A methodology to pipeline combinational circuits

1. Draw a line crossing all outputs
2. Draw a line to partition the graph into two parts;
 - all the arrows must cross the partition line in the same direction
 - add a pipeline register at every point where the partition line crosses an edge


$$P1: t_{CLK} = \max\{4+3+8, 4+4, 4+4+2, 5\} = 15$$

$$P2: t_{CLK} = \max\{4+3, 8+5, 4+5, 4+2+5\} = 13$$

3. To increase the number of pipeline stages, further partition any partition using step 2

$$P1+P2: t_{CLK} = \max\{4+3, 8, 4+2, 5\} = 8$$

4. For the fastest possible pipeline make the slowest box be the determinant of the clock speed, i.e., the bottleneck



Homework:

Draw another three stage pipeline that does not increase the clock period


Latency and throughput of pipelined combinational circuit

Clock Period = 8 ns

Latency through full circuit:

$$3 \times 8 = 24 \text{ ns}$$

Throughput = $1/8 \text{ ns}$




Tools for pipelining combinational circuits


- Once we specify the number of stages or the clock period, modern tools can pipeline combinational circuits to generate fairly optimal pipelines
 - The method tools follow is similar to the one shown on the previous slide except that the tools work with networks of primitive gates

We won't be using such tools in this class!

Elastic vs Inelastic pipeline



Inelastic: Uses registers to separate pipeline stages;
All pipeline stages move synchronously;
It is complicated to start and stop a synchronous pipeline



Elastic: Uses Fifos to separate pipeline stages;
A stage can process data if its input FIFO is not empty
and output FIFO is not Full

In this course we will focus exclusively on Elastic pipelines
Why?

Inelastic vs Elastic Pipelining

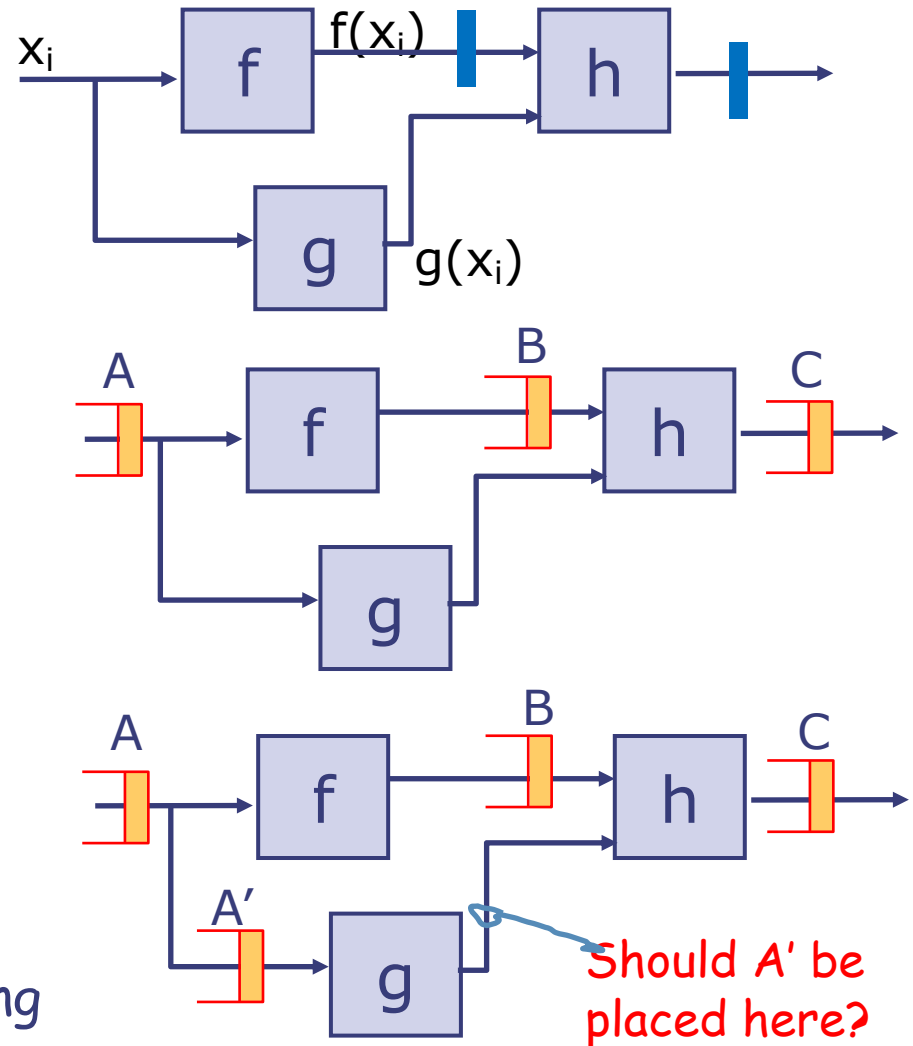
Incorrect synchronous pipeline
output = $h(f(x_i), g(x_{i+1}))$

Assuming fifo B is empty initially, a rule describing this circuit will never fire

Introduction of an extra fifo
A' will remove the deadlock

- rule 1 will enq in A' and B
- rule 2 will enq in C

It is easier to preserve the correctness using fifos but balancing is still an issue




Pipelining sequential modules

- Now suppose each box is sequential circuit (aka module) instead of a combinational circuit
- Let the number in each module represent the number of clock cycles that module takes to produce the output given an input
- Pipelined circuit still makes sense

Stage S1: 7 cycles

Stage S2: 8 cycles

Stage S3: 5 cycles




However, on an average we can't process inputs any faster than the slowest stage. After some time the output fifo of fast stage will get blocked by following slow stage

Cascaded Fifos



- When we draw a fifo, we usually don't indicate the number of elements in it.
- Thus, cascaded fifo's can be represented by one fifo
- Also for analysis purpose we will assume that the input Fifos of the system are never empty and output Fifos are never full

Pipeline bubbles



- Suppose f takes 1 or 2 cycles depending upon the input data, and g always takes one cycle
- Suppose input data to this system are x_0, x_1, x_2, x_3 , where x_0, x_2 take one cycle each and x_1, x_3 take 2 cycles each



	t0	t1	t2	t3	t4	t5	t6	t7	t8
f	x_0	x_1 — x_1	x_2	x_3 — x_3					
g		x_0		x_1	x_2		x_3		

Diagram illustrating the pipeline execution over time (t0 to t8). The pipeline has two stages, f and g. Red arrows show the data flow: x_0 flows from f at t0 to g at t1; x_1 flows from f at t1 to g at t3; x_2 flows from f at t2 to g at t4; x_3 flows from f at t3 to g at t6. Blue clouds represent pipeline bubbles: one in stage g at t2, and one in stage f at t5. A line points from the text 'A bubble is introduced any time a pipeline stage has no work' to the bubble in stage g at t2.

A bubble is introduced any time a pipeline stage has no work

A bubble represents a loss of performance

Feed-Forward Pipelines

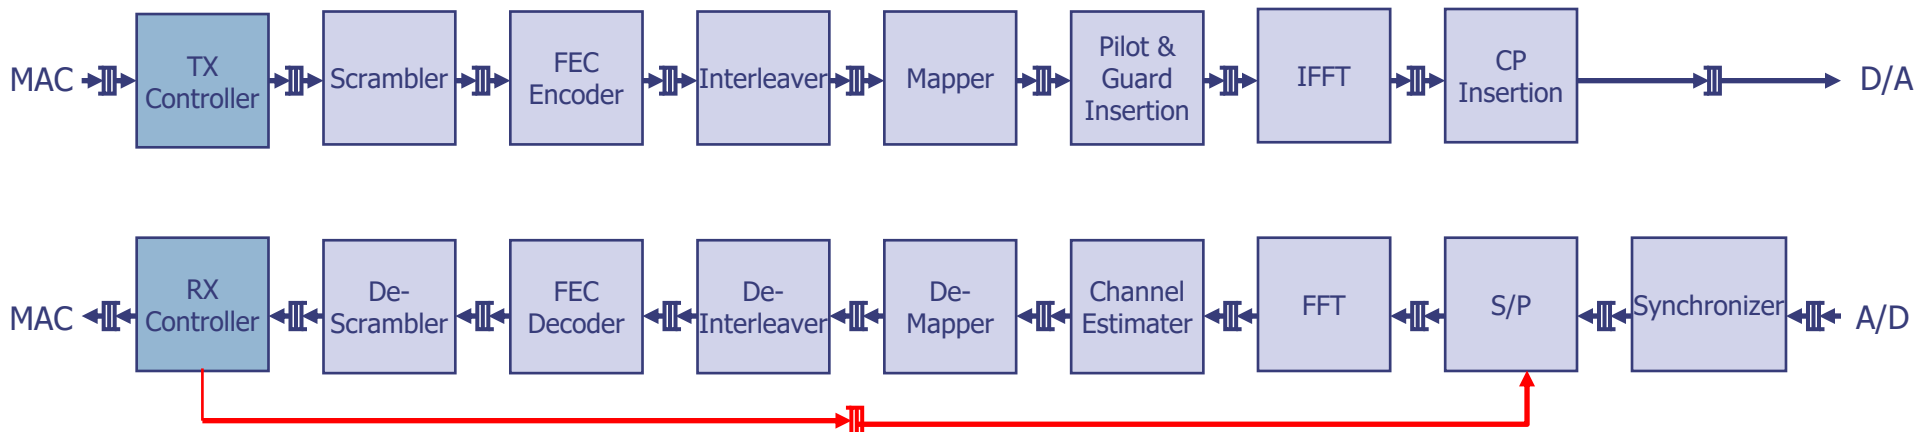


- Feed-forward pipelines have no feedback loops
 - Easy to insert pipeline stages and preserve functionality
- One can easily alter the functionality of a pipeline with feedback loops by inserting a fifo in the feed back path
- Pipelining a network of modules with loops requires careful analysis of what the pipeline is supposed to do
- Analysis is difficult because a module's interface does not say
 - if the module is internally pipelined
 - if it processes items in order

Some examples

Example of feed-forward pipeline

Baseband OFDM protocols



- WiFi, WiMax, WUSB are all examples of OFDM protocols
- All stages behave strictly in the first-in first-out manner
- Stages differ enormously in terms of complexity and area
- In spite of the similar structure, specific protocols have very different performance requirements, e.g.,
 - WiFi: 64pt @ 0.25MHz
 - WiMAX: 256pt @ 0.03MHz
 - WUSB: 128pt 8MHz
- There is a feedback path but it is only to shut down the pipeline

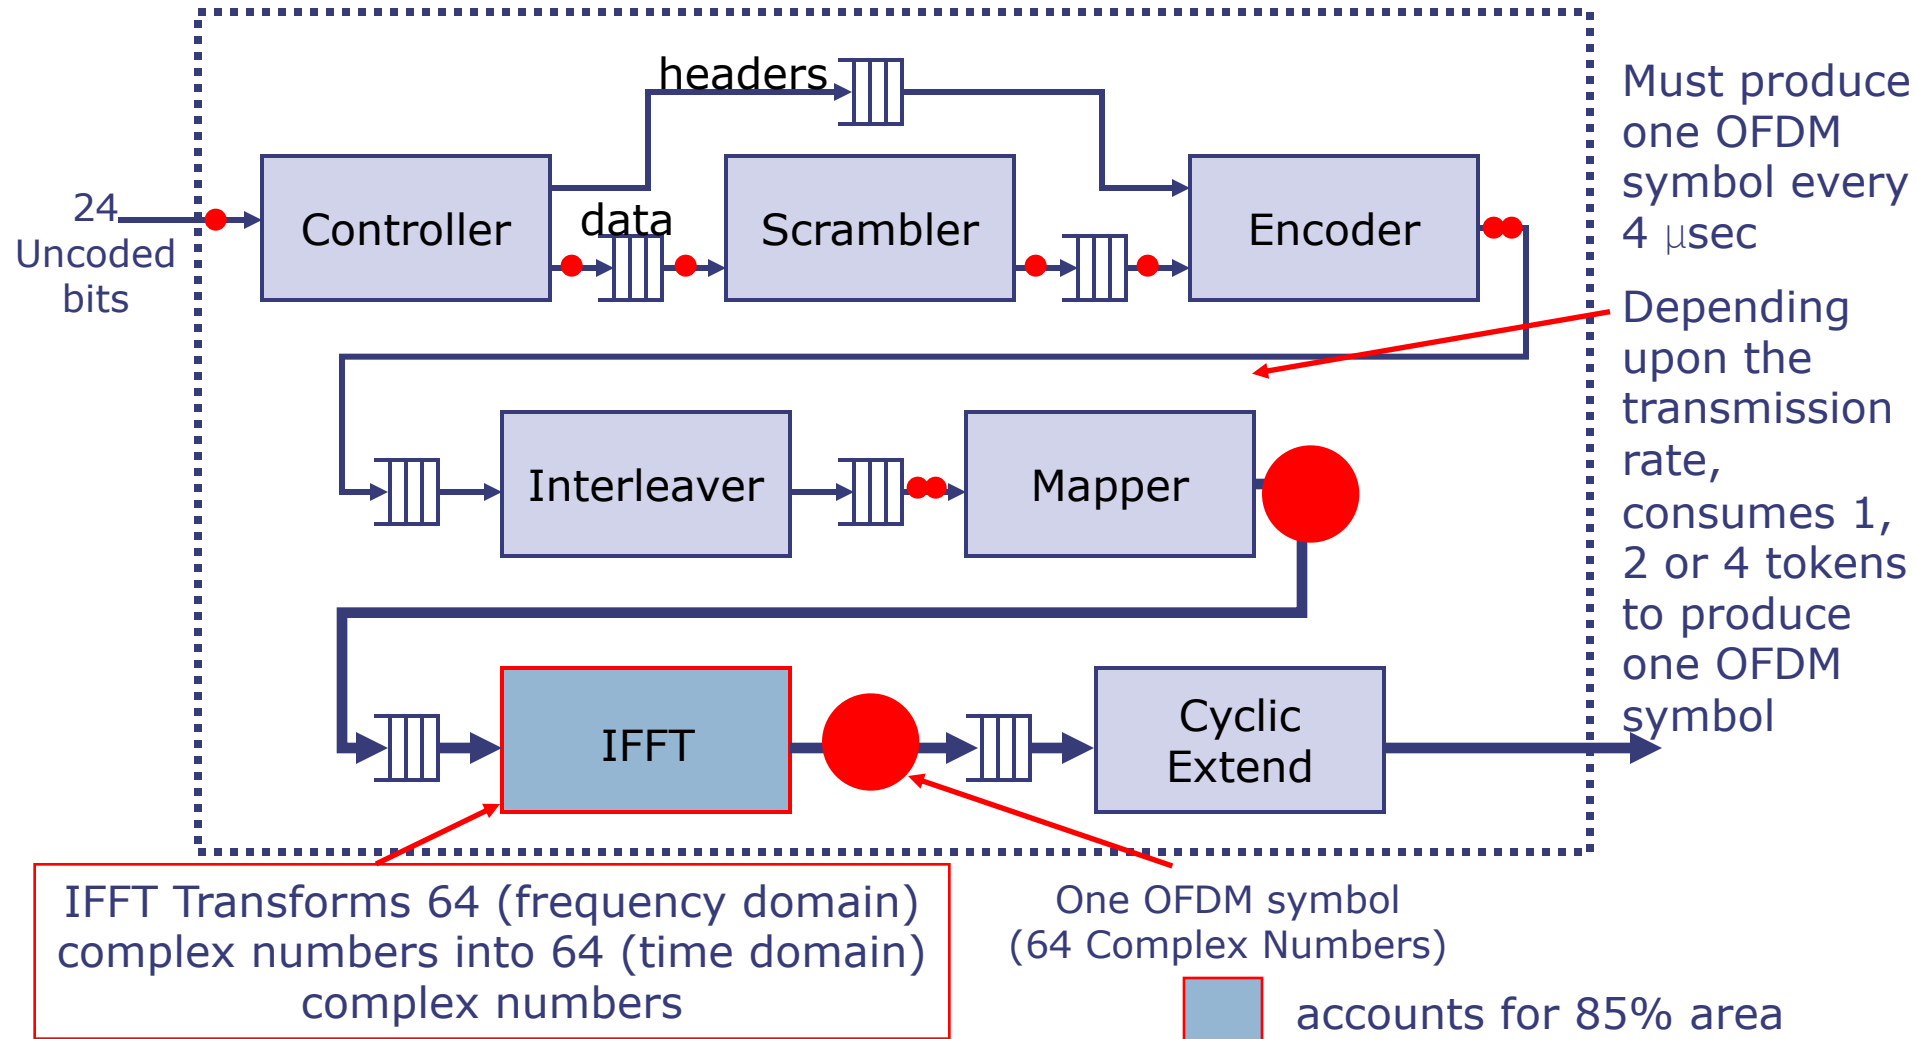
no brownie points for
running the protocol
faster than the specs

Designing feed-forward pipelines

- Requires optimizing area and clock speed to consume least amount of power while meeting the performance requirements
 - Some phases are much more complex computationally than others, and may have to be pipelined internally to meet the performance goals
 - Some phases can be turned into multicycle implementations to reduce area without affecting the performance
 - Some simple phases can be merged to save area

Design exploration is essential for good designs;
Modules with Latency-Insensitive interfaces
facilitate design exploration

802.11a Transmitter



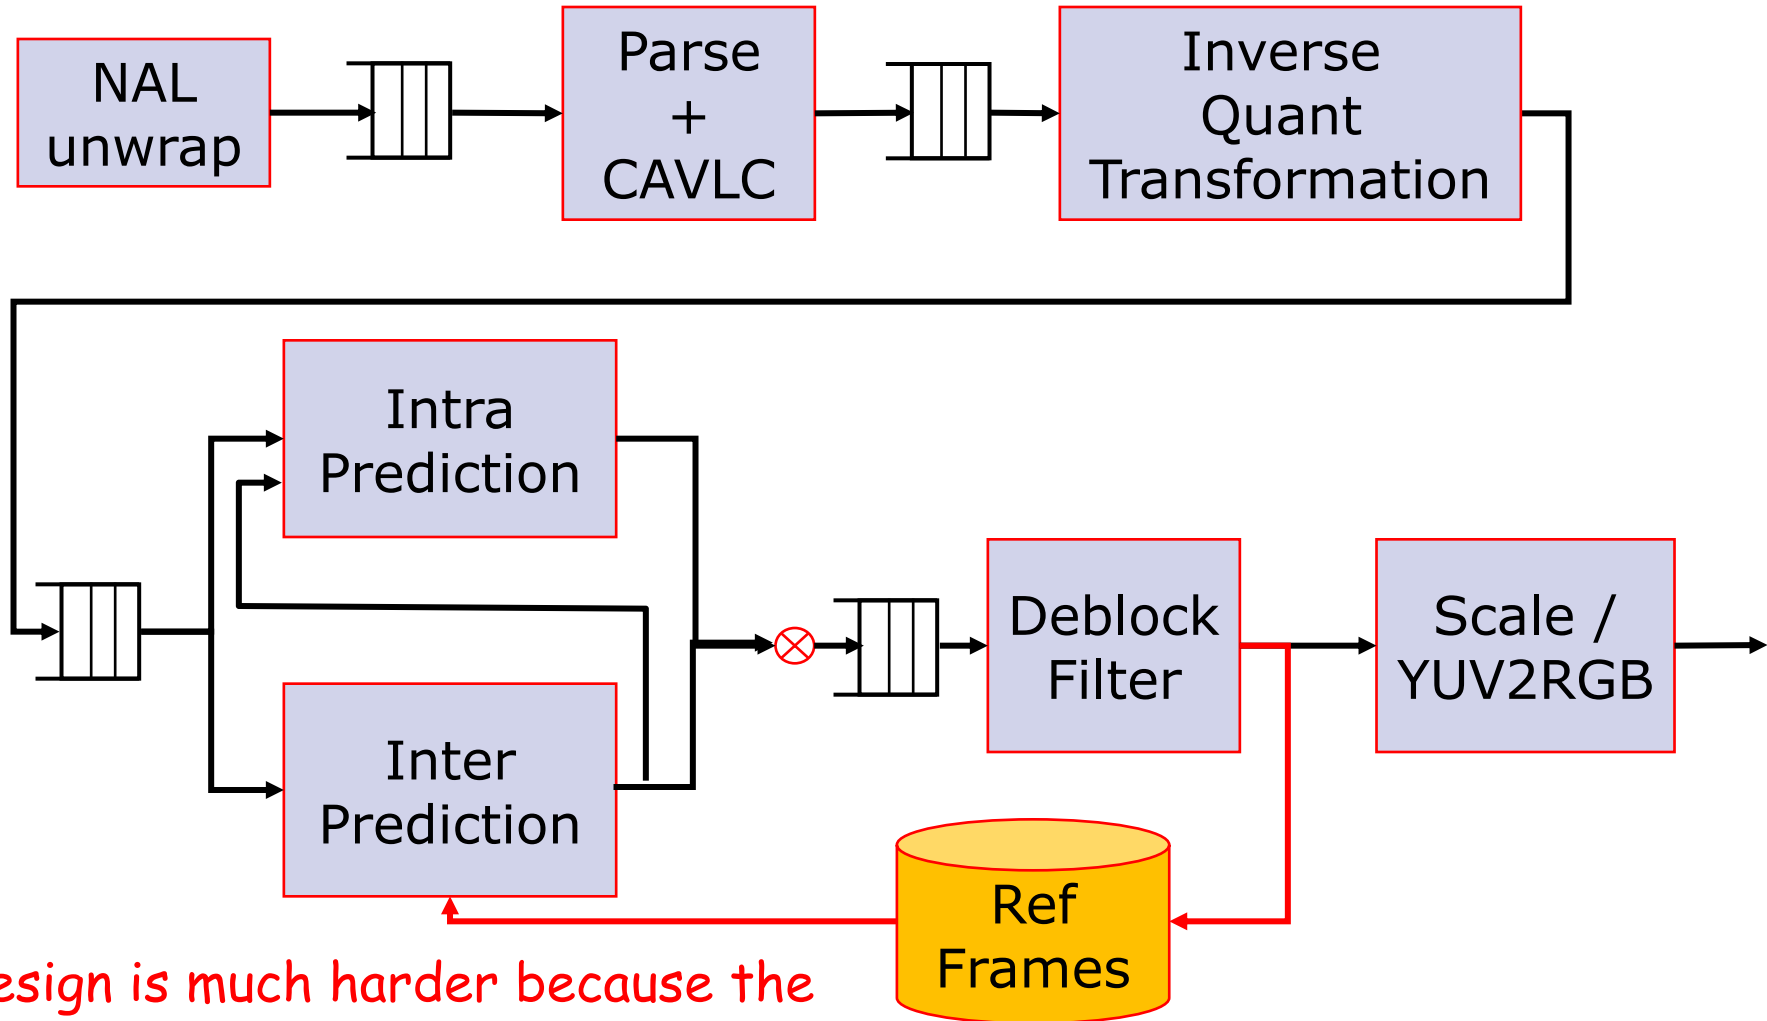
802.11a Transmitter Design:

Design Block	Lines of Code (BSV)	Relative Area
Controller	49	0%
Scrambler	40	0%
Conv. Encoder	113	0%
Interleaver	76	1%
Mapper	112	11%
IFFT	95	85%
Cyc. Extender	23	3%

Complex arithmetic libraries constitute another 200 lines of code

Pipelines with feedback

H.264 Decoder



Design is much harder because the functionality depends upon the feedback loop

Processor pipelines

- Pipelining processor provides the ultimate challenge in computer architecture
 - Stringent correctness requirements
 - Requires speculative execution of instructions to pipeline at all!
 - Requires dealing with a variety of feedbacks
 - The goal is always to achieve highest performance but within a given area and power budget

Next two lectures and your project will be on processor pipelines