The Memory Hierarchy

Reminders:
Lab 4 due April 7
Lab 3 checkoff due April 8
Quiz 2 Review: Tuesday, April 14, 7:30pm EST
Quiz 2: Thursday, April 16 7:30pm EST

Memory Technologies

- Technologies have vastly different tradeoffs between capacity, latency, bandwidth, energy, and cost
  - … and thus different applications

	Capacity	Latency	Cost/GB	
Register	100s of bits	20 ps	$$$$	**Processor Datapath**
SRAM	~10 KB-10 MB	1-10 ns	~$1000	**Memory Hierarchy**
DRAM	~10 GB	80 ns	~$10	
Flash*	~100 GB	100 us	~$1	**I/O subsystem**
Hard disk*	~1 TB	10 ms	~$0.1	

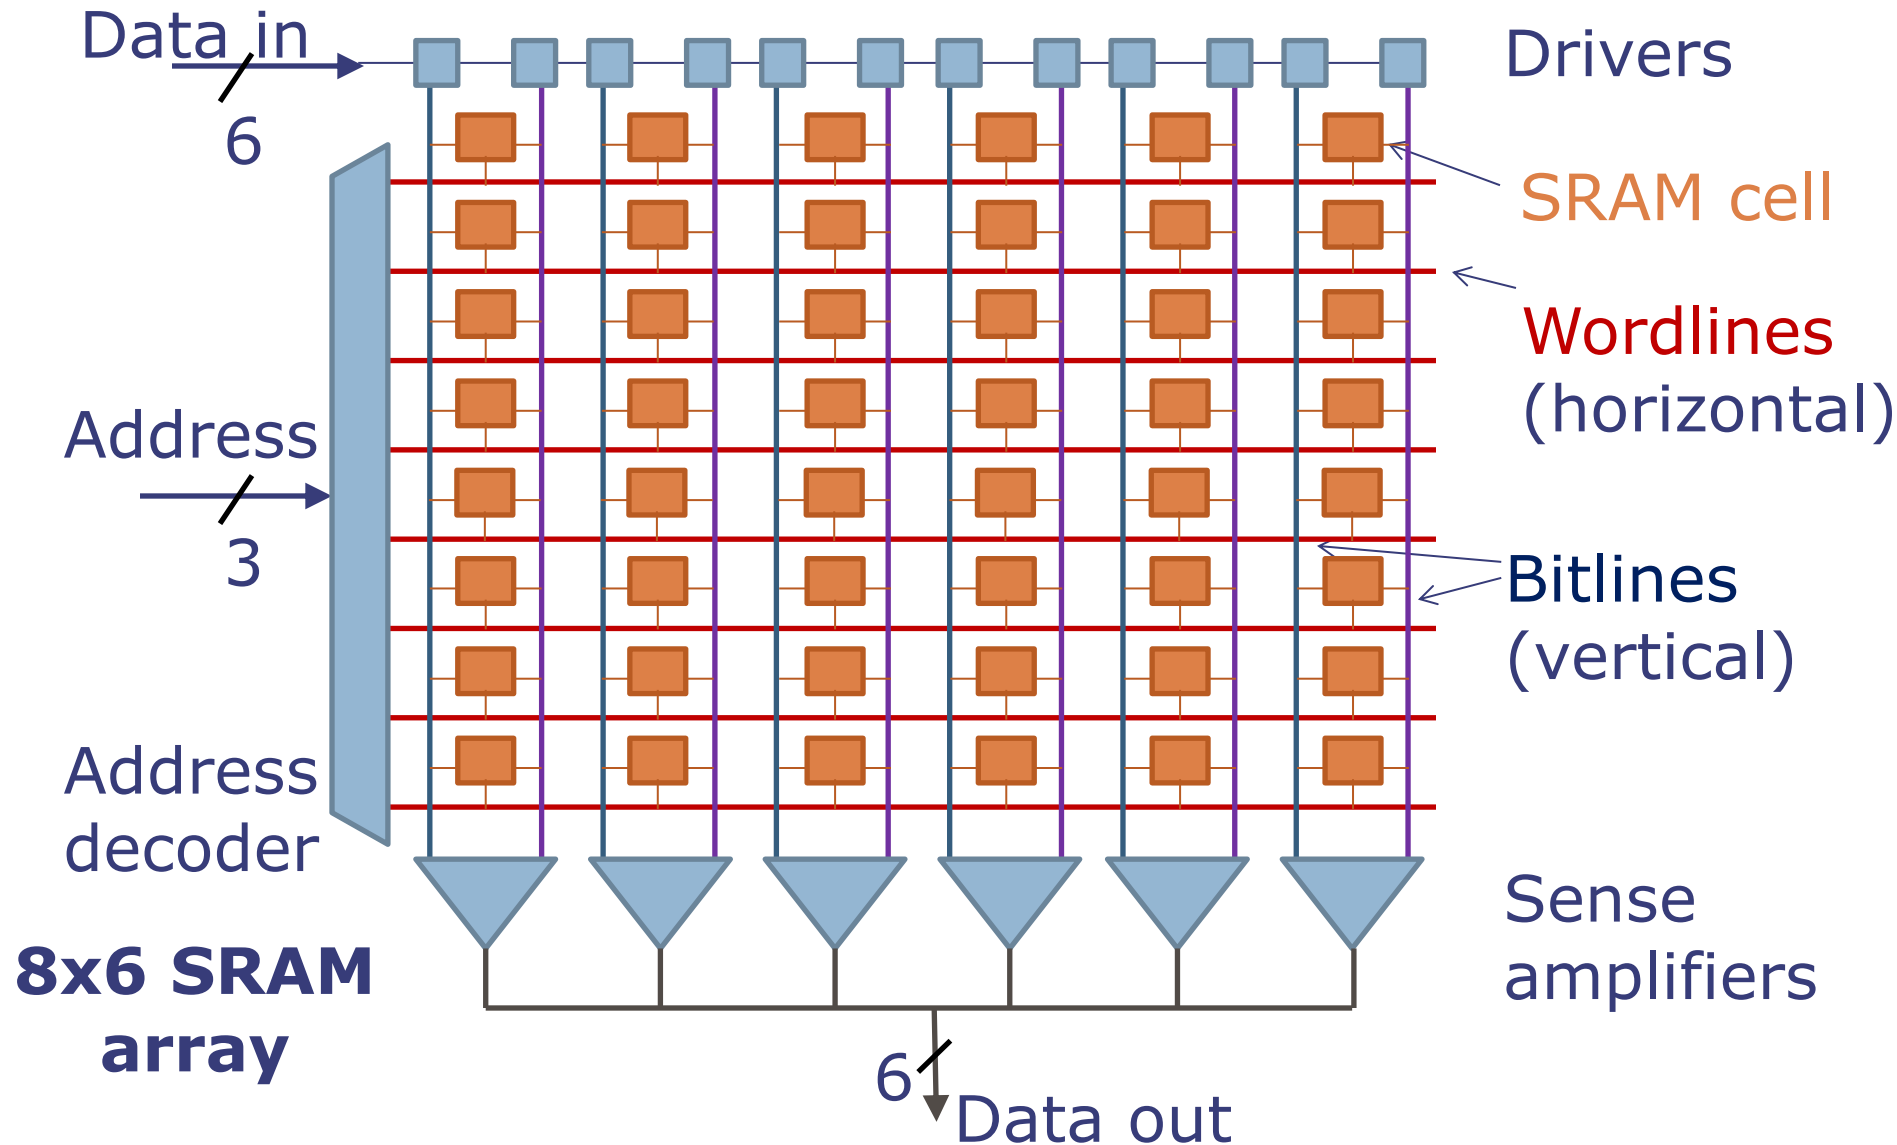* non-volatile (retains contents when powered off)

Memory Technologies: SRAM, DRAM, Flash, Hard Disk

## *NOTE: Demystification, will not be on the quiz*

Static RAM (SRAM)

Data in

6

Address

3

Address
decoder

**8x6 SRAM
array**

Drivers

SRAM cell

Wordlines
(horizontal)

Bitlines
(vertical)

Sense
amplifiers

6

Data out

SRAM Cell

6-transistor (6T) cell:

- Two CMOS inverters (4 FETs) forming a bistable element
- Two access transistors

bitline $\overline{\text{bitline}}$

6T SRAM Cell

Wordline

access FETs

Bistable element
(two stable states)
stores a single bit
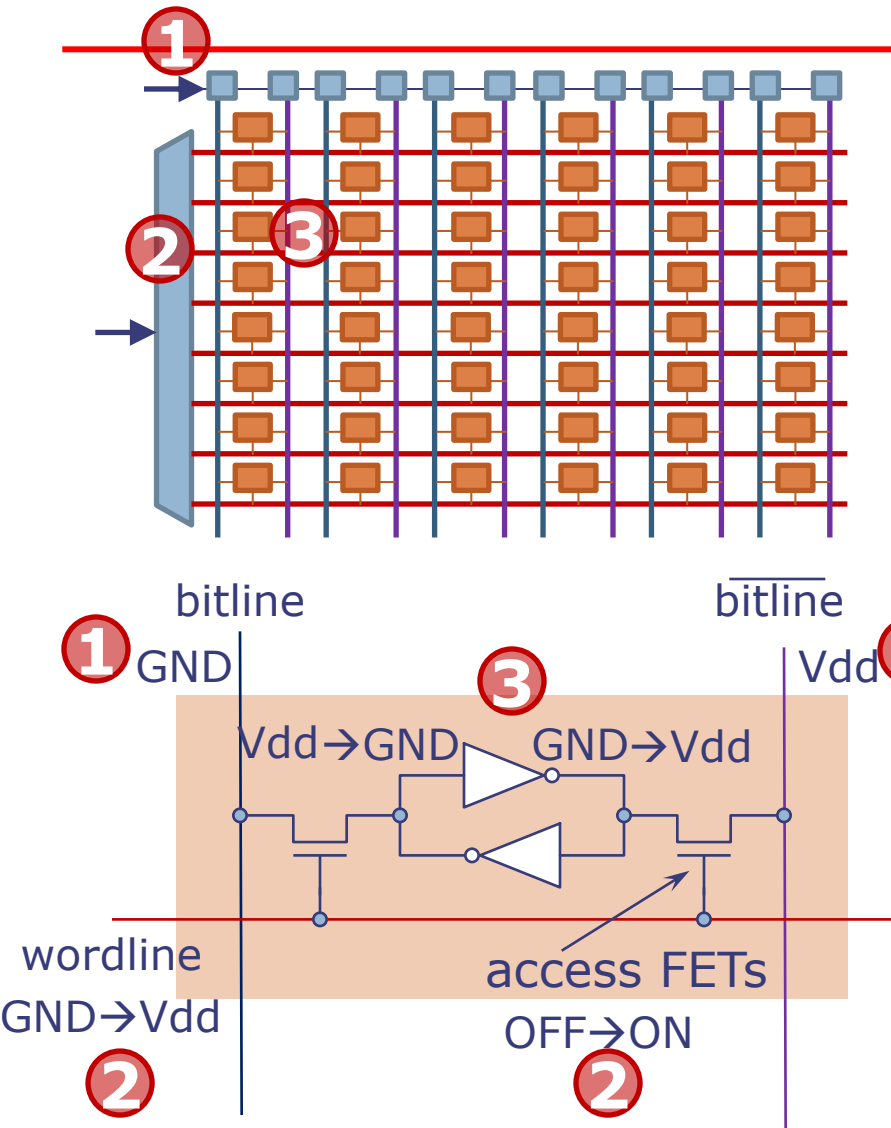
Vdd — GND "1"

GND — Vdd "0"

SRAM Read

1. **Drivers precharge all bitlines to Vdd (1), and leave them floating**

2. **Address decoder activates one wordline**

3. **Each cell in the activated word slowly pulls down one of the bitlines to GND (0)**

4. **Sense amplifiers sense change in bitline voltages, produce output data**

bitline $\overline{\text{bitline}}$

Vdd   6T SRAM Cell   Vdd

1   0

access FETs

wordline
GND→Vdd   V(t)   OFF→ON   V(t)

t   t

SRAM Write

1. Drivers set and hold bitlines to desired values (Vdd and GND for 1, GND and Vdd for 0)

2. Address decoder activates one wordline

3. Each cell in word is overpowered by the drivers, stores value

Cell transistors are carefully sized so that bitline GND overpowers cell Vdd, but bitline Vdd does not overpower cell GND

bitline

$\overline{\text{bitline}}$

GND

Vdd→GND   GND→Vdd

Vdd

access FETs

wordline

GND→Vdd

OFF→ON

Multiported SRAMs

- SRAM so far can do either one read or one write/cycle

- We can do multiple reads and writes with multiple ports by adding one set of wordlines and bitlines per port

- *Cost/bit for N ports?*
  - *Wordlines?* N
  - *Bitlines?* 2*N
 - *Access FETs?* 2*N

- Wires dominate area → $O(N^2)$ area!

Summary: SRAMs

- Array of k*b cells (k words, b cells per word)
- Cell is a bistable element + access transistors
  - Analog circuit with carefully sized transistors
- Read: Precharge bitlines, activate wordline, sense
- Write: Drive bitlines, activate wordline, overpower cells

- 6 FETs/cell… can we do better?

1T Dynamic RAM (DRAM) Cell

Cyferz (CC BY 2.5)

Storage
capacitor

1T DRAM Cell

wordline

access FET

V_{REF}
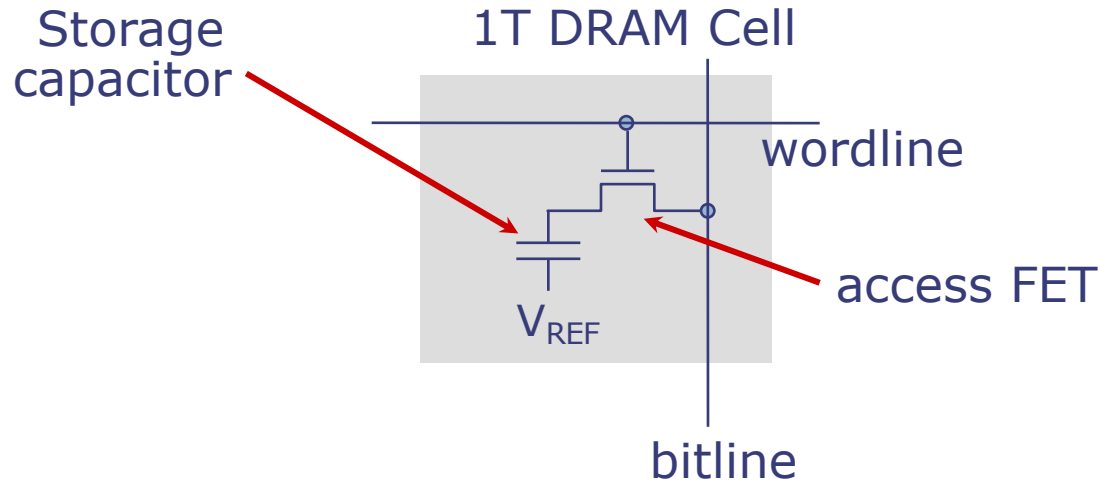
bitline

Trench capacitors
take little area

✓ ~20x smaller area than SRAM cell → Denser and cheaper!
✗ Problem: Capacitor leaks charge, must be refreshed periodically (~milliseconds)

DRAM Writes and Reads

- Writes: Drive bitline to Vdd or GND, activate wordline, charge or discharge capacitor

1T DRAM Cell

word line

Storage capacitor

V_{REF}

access FET

bitline

- Reads:
  1. Precharge bitline to Vdd/2
  2. Activate wordline
  3. Capacitor and bitline share charge
     - If capacitor was discharged, bitline voltage decreases slightly
     - If capacitor was charged, bitline voltage increases slightly
  4. Sense bitline to determine if 0 or 1

- Issue: Reads are destructive! (charge is gone!)
 - Data must be rewritten to cells at end of read
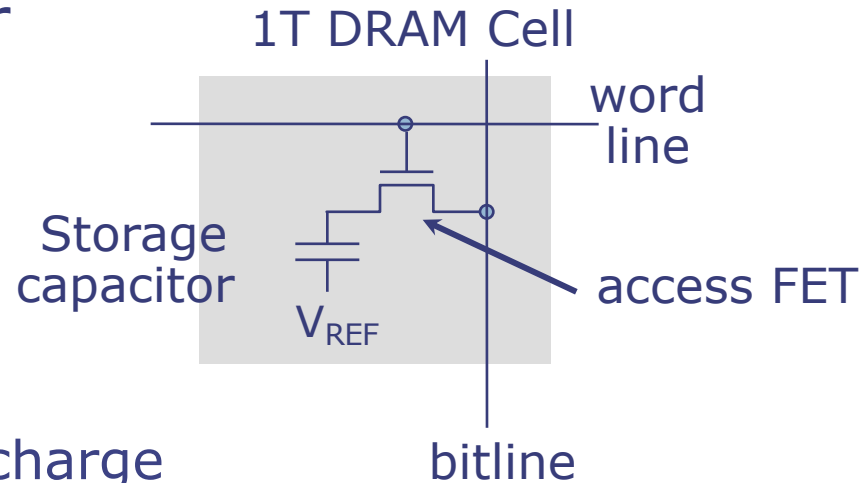
Non-Volatile Storage: Flash
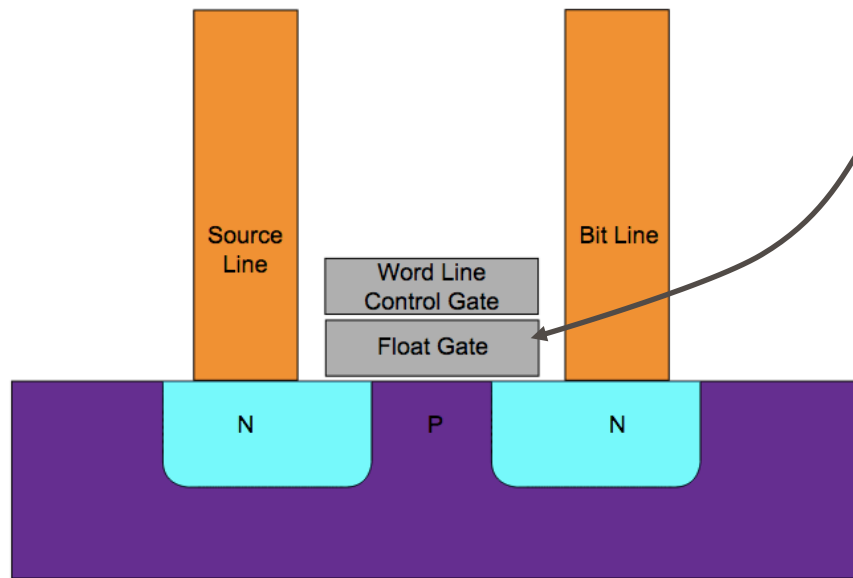
Electrons here diminish strength of field from control gate ⇒ no inversion ⇒ NFET stays off even when word line is high.

Cyferz (CC BY 2.5)

Flash Memory: Use "floating gate" transistors to store charge (floating gate is a well insulated conductor)
- Very dense: Multiple bits/transistor, read and written in blocks
- Slow (especially on writes), 10-100 us
- Limited number of writes: charging/discharging the floating gate (writes) requires large voltages that damage transistor

Non-Volatile Storage: Hard Disk

Platter
Spindle
Head
Actuator Arm
Actuator Axis
Power Connector
Jumper Block
IDE Connector
Actuator

Surachit (CC BY 2.5)

Disk head

Circular track
divided into sectors

Hard Disk: Rotating magnetic platters + read/write head
- Extremely slow (~10ms): Mechanically move head to position, wait for data to pass underneath head
- ~100MB/s for sequential read/writes
- ~100KB/s for random read/writes
- Cheap

The Memory Hierarchy

Summary: Memory Technologies

	Capacity	Latency	Cost/GB
Register	100s of bits	20 ps	$$$$
SRAM	~10 KB-10 MB	1-10 ns	~$1000
DRAM	~10 GB	80 ns	~$10
Flash	~100 GB	100 us	~$1
Hard disk	~1 TB	10 ms	~$0.1

- **Different technologies have vastly different tradeoffs**
- **Size is a fundamental limit, even setting cost aside:**
  - Small + low latency, high bandwidth, low energy, or
  - Large + high-latency, low bandwidth, high energy
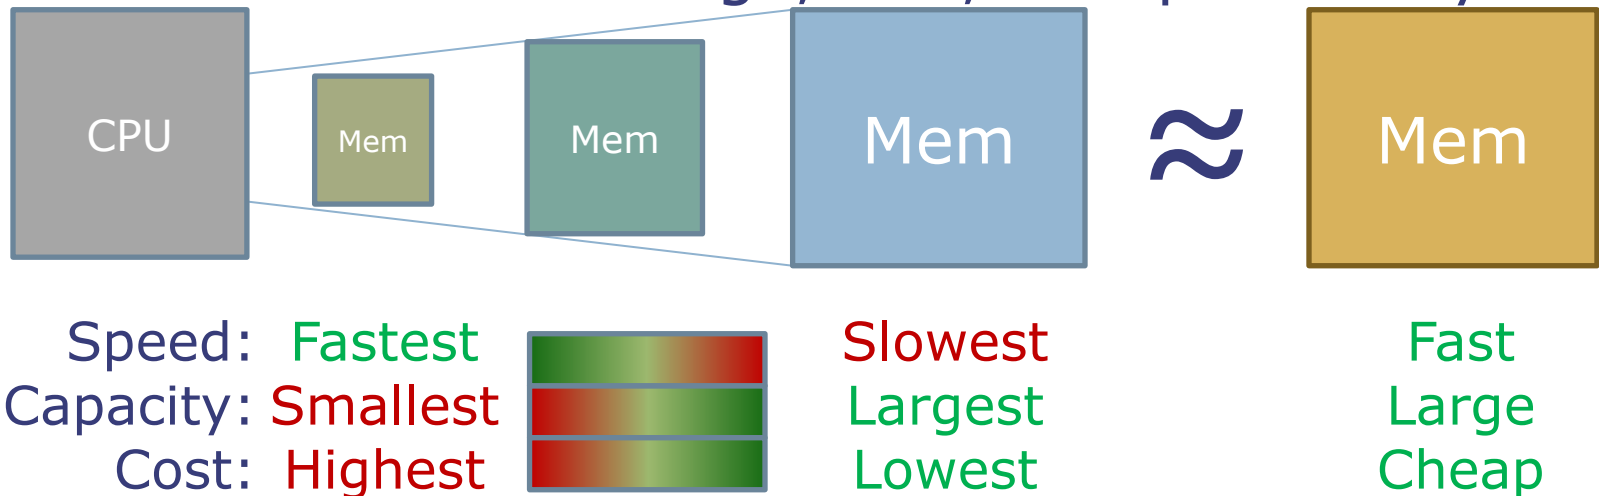- **Can we get best of both worlds? (large, fast, cheap)**

The Memory Hierarchy

Want large, fast, and cheap memory, but...

Large memories are slow (even if built with fast components)

Fast memories are expensive

Solution: Use a hierarchy of memories with different tradeoffs to fake a large, fast, cheap memory

	CPU	Mem	Mem	Mem	≈	Mem

Speed:	Fastest	Slowest	Fast
Capacity:	Smallest	Largest	Large
Cost:	Highest	Lowest	Cheap

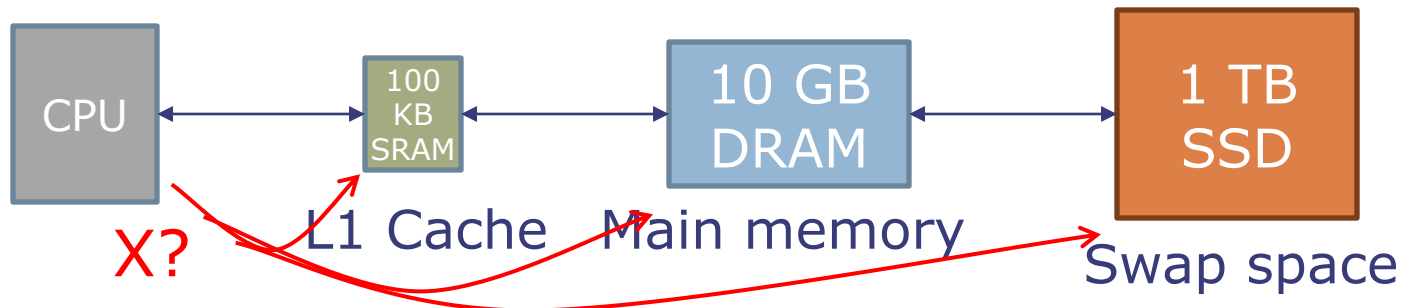# Memory Hierarchy Interface

Approach 1: Expose Hierarchy

- Registers, SRAM, DRAM, Flash, Hard Disk each available as storage alternatives

- Tell programmers: "Use them cleverly"

Approach 2: Hide Hierarchy

- Programming model: Single memory, single address space
- Machine transparently stores data in fast or slow memory, depending on usage patterns

X?   L1 Cache   Main memory   Swap space

Typical Memory Access Patterns

array accesses

address

data

local variable accesses

procedure calls

stack

loop

code

time

Common Predictable Patterns

- Two predictable properties of memory accesses:

  - Temporal locality: If a location has been accessed recently, it is likely to be accessed (reused) soon

  - Spatial locality: If a location has been accessed recently, it is likely that nearby locations will be accessed soon

Caches

- Cache: A small, interim storage component that transparently retains (caches) data from recently accessed locations

- Processor sends accesses to cache. Two options:
  - Cache hit: Data for this address in cache, returned quickly
  - Cache miss: Data not in cache
    - Fetch data from memory, send it back to processor
    - Retain this data in the cache (replacing some other data)
  - Processor must deal with variable memory access time

A Typical Memory Hierarchy

Computers use many levels of caches:

On the CPU

On chip

Other chips and devices

Registers
Level 1 Cache
Level 2 Cache
Level 3 Cache
Main Memory
Flash Drive
Hard Disk

	Access time	Capacity	Managed By
Registers	1 cycle	1 KB	Software/Compiler
Level 1 Cache	2-4 cycles	32 KB	Hardware
Level 2 Cache	10 cycles	256 KB	Hardware
Level 3 Cache	40 cycles	10 MB	Hardware
Main Memory	200 cycles	10 GB	Software/OS
Flash Drive	10-100us	100 GB	Software/OS
Hard Disk	10ms	1 TB	Software/OS

Cache Metrics

- Hit Ratio: $HR = \dfrac{hits}{hits + misses} = 1 - MR$

- Miss Ratio: $MR = \dfrac{misses}{hits + misses} = 1 - HR$

- Average Memory Access Time (AMAT):

$$AMAT = HitTime + MissRatio \times MissPenalty$$

- Goal of caching is to improve AMAT
- Formula can be applied recursively in multi-level hierarchies:

$$AMAT = HitTime_{L1} + MissRatio_{L1} \times AMAT_{L2} =$$

$$AMAT = HitTime_{L1} + MissRatio_{L1} \times (HitTime_{L2} + MissRatio_{L2} \times AMAT_{L3}) = \dots$$

Example: How High of a Hit Ratio?

CPU ←→ Cache ←→ Main Memory

4 cycles   100 cycles

What hit ratio do we need to break even?
(Main memory only: AMAT = 100)

$$100 = 4 + (1 - HR) \times 100 \Rightarrow HR = 4\%$$

What hit ratio do we need to achieve AMAT = 5 cycles?

$$5 = 4 + (1 - HR) \times 100 \Rightarrow HR = 99\%$$

Basic Cache Algorithm (Reads)

CPU

Tag   Data

A	Mem[A]
B	Mem[B]

$(1-HR)$

Main Memory

On reference to Mem[X],
look for X among cache tags

HIT: X = Tag(i)
for some
cache line i

Return Data(i)

MISS: X not
found in Tag
of any cache line

Read Mem[X]
Return Mem[X]
Select a line k
to hold Mem[X]
Write Tag(k)=X,
Data(k) = Mem[X]

Q: How do we "search" the cache?

Thank you!

Next lecture: Cache Tradeoffs